

VDMJ Design Specification	
Author	Nick Battle
Date	16/04/10
Issue	1.0

0. Document Control

0.1. Table of Contents

0. Document Control.....	2
0.1. Table of Contents.....	2
0.2. References.....	2
0.3. Document History.....	3
0.4. Copyright.....	3
1. Overview.....	4
1.1. Package Overview.....	4
2. Package Detail.....	6
2.1. vdmj.....	6
2.2. vdmj.lex.....	7
2.3. vdmj.syntax.....	10
2.4. vdmj.ast.....	12
2.5. vdmj.types.....	12
2.6. vdmj.expressions.....	14
2.7. vdmj.statements	16
2.8. vdmj.patterns.....	17
2.9. vdmj.traces.....	18
2.10. vdmj.definitions.....	20
2.11. vdmj.modules.....	23
2.12. vdmj.typechecker.....	24
2.13. vdmj.pog.....	26
2.14. vdmj.runtime.....	29
2.15. vdmj.scheduler.....	34
2.16. vdmj.values.....	44
2.17. vdmj.commands.....	47
2.18. vdmj.messages.....	47
2.19. vdmj.debug.....	48
2.20. vdmj.config.....	51
2.21. vdmj.util.....	52
3. External Interfaces.....	53
3.1. The Native Call Interface.....	53
3.2. The Remote Control Interface.....	53

0.2. References

- [1] Wikipedia entry for The Vienna Development Method,
 http://en.wikipedia.org/wiki/Vienna_Development_Method

- [2] Wikipedia entry for Specification Languages, http://en.wikipedia.org/wiki/Specification_language
- [3] The VDM Portal, <http://www.vdmportal.org/twiki/bin/view>
- [4] The VDMTools VDM-SL Language Manual, http://www.vdmtools.jp/uploads/manuals/langmansl_a4E.pdf
- [5] The VDMTools VDM++ Language Manual, http://www.vdmtools.jp/uploads/manuals/langmanpp_a4E.pdf
- [6] DBGp - A common debugger protocol for languages and debugger UI communication, <http://xdebug.org/docs-dbgp.php>.
- [7] Overture - Open-source Tools for Formal Modelling, <http://www.overturetool.org/>.
- [8] Modelling and Validating Distributed Embedded Real-Time Control Systems, Marcel Verhoef, PhD Thesis.

0.3. Document History

Issue 0.1	22/10/08	First release.
Issue 0.2	28/11/08	Added comments from PGL. Added AST converter.
Issue 0.3	04/03/09	Added PO generator section, vdm.traces and misc other changes.
Issue 0.4	28/05/09	Added the DBGp protocol section, and extended runtime to discuss class initialization.
Issue 0.5	17/09/09	Added detail about VDM-RT implementation.
Issue 0.6	02/10/09	Updated CT description for TraceVariables
Issue 0.7	09/12/09	Added GPL copyright section.
Issue 0.8	22/01/10	Added section 3. Minor updates.
Issue 0.9	18/02/10	Various updates, and added section 4.
Issue 1.0	16/04/10	Updated for Overture 0.2.0, added new scheduling details

0.4. Copyright

Copyright © 2010, Fujitsu Services Ltd.

Permission is granted to copy, distribute and/or modify this document under the terms of the GNU Free Documentation License, Version 1.3 or any later version published by the Free Software Foundation; with no Invariant Sections, no Front-Cover Texts, and no Back-Cover Texts.

1. Overview

VDMJ provides tool support for the VDM-SL, VDM++ and VDM-RT specification languages, written in Java [4][5][8]. The tool includes a parser, a type checker, an interpreter, a debugger and a proof obligation generator. It is a command line tool only, though it is accessible from graphical environments like Eclipse [7].

1.1. Package Overview

The implementation is divided into the following Java packages, which are all sub-packages of `org.overturetool`.

Packages	
<code>vdmj</code>	The main VDMJ class and supporting classes.
<code>vdmj.lex</code>	Classes that implement the lexical analyser and its tokens.
<code>vdmj.syntax</code>	Classes that implement the syntax analyser.
<code>vdmj.ast</code>	Classes that translate an Overture AST to VDMJ's internal tree.
<code>vdmj.types</code>	Classes that represent static types during type checking.
<code>vdmj.expressions</code>	Classes that represent VDM expressions.
<code>vdmj.statements</code>	Classes that represent operation statements.
<code>vdmj.patterns</code>	Classes that represent patterns and binds.
<code>vdmj.traces</code>	Classes that represent trace definitions.
<code>vdmj.definitions</code>	Classes that represent VDM definitions.
<code>vdmj.modules</code>	Classes that represent VDM-SL modules and their import/export definitions.
<code>vdmj.typechecker</code>	Classes that support the static type checker.
<code>vdm.pog</code>	Classes that support the proof obligation generator.
<code>vdmj.runtime</code>	Classes that implement the interpreter.
<code>vdmj.scheduler</code>	Classes that implement the deterministic thread scheduler.
<code>vdmj.values</code>	Classes that represent runtime values in the interpreter.
<code>vdmj.commands</code>	Classes that read and execute commands from standard input.
<code>vdmj.messages</code>	Classes that hold VDMJ error and warning messages.
<code>vdmj.debug</code>	Classes that implement the DBGp protocol.
<code>vdmj.config</code>	Classes that load and access property file settings.
<code>vdmj.util</code>	Utility classes and library routines used by all other packages.

The `vdmj` package contains the “main” abstract class for the suite, with subclasses to parse, type check and interpret a specification in the VDM-SL, VDM++ or VDM-RT dialects.

The `vdmj.lex` package contains all classes to do with the lexical analysis of specifications. This includes the lexical token reader, plus a set of value classes for representing the various types of lexical token.

The `vdmj.syntax` package contains all the syntax analysis classes. This includes eight “readers”, which implement a backtracking recursive descent parser, based on a stream of lexical tokens. The readers are all sub-classes of an abstract `Reader` class.

The `vdmj.ast` package contains classes to translate an Overture parsed AST into VDMJ's internal tree

format. This was originally to permit the Overture AST to be used by VDMJ, but without changing the type checking and runtime classes. It is no longer used.

The `vdmj.types` package contains value classes that represent the various static types that can be contained in a VDM specification. They are all sub-classes of an abstract `Type` class.

The `vdmj.expressions` package contains value classes that represent the various types of expression that can be defined in a specification. They are all sub-classes of an abstract `Expression` class, which defines methods for an expression to be type checked and evaluated. Similarly, the `vdmj.statements` package defines a set of value classes that subclass `Statement` and represent the different statements in a specification.

The `vdmj.patterns` package includes value classes to represent the various patterns and binds in a specification. They are all subclasses of an abstract `Pattern` or `Bind` class.

The `vdmj.traces` package includes classes that represent the possible trace definitions that can be declared in a VDM-SL, VDM++ or VDM-RT class.

The `vdmj.definitions` package contains a set of classes representing the definitions in a specification. They are all sub-classes of an abstract `Definition` class. Definitions are nested, so for example a class definition may contain function definitions, which in turn may contain function definitions for their pre and post condition functions. All definitions implement methods to perform type checking, and to generate runtime values representing their content.

The `vdmj.modules` package contains value classes that represent the modular structure of VDM-SL specifications, including their import and export declarations.

The `vdmj.typechecker` package includes classes which control the type checking of specifications. Most of the type checking is performed by the definition, expression and statement classes that collectively describe the specification, but the typechecker package defines the supporting classes to invoke the type checking methods of the other objects, and to represent the static environment in which type checking is performed.

The `vdmj.pog` package contains classes that support the proof obligation generator. Like type checking, the actual process of proof obligation generation is performed by the definitions, expressions and statements which form the specification, but the `pog` package contains classes to represent proof obligations.

The `vdmj.runtime` package defines the abstract interpreter, and its subclasses to interpret specifications. It includes classes to represent the runtime execution context, and exceptions which can be generated at runtime.

The `vdmj.scheduler` package defines classes that schedule the deterministic execution of multiple threads, and for VDM-RT, coordinate the movement of time.

The `vdmj.values` class contains a set of classes to represent runtime values in the interpretation of a specification. They are all subclasses of the abstract `Value` class. All values are immutable, with the exception of the `UpdatableValue` class hierarchy, which has a `set` method and is used to implement all state variables in the system.

The `vdmj.commands` package contains the command line readers that implement the interactive actions of the suite. This should be the only package that interacts with the user terminal.

The `vdmj.messages` package contains value classes and exceptions for holding error and warning messages.

The `vdmj.debug` package contains classes which implement the DBGp protocol described in [6]. This is used by the Eclipse debugger in Overture to debug specifications using VDMJ.

The `vdmj.config` package contains classes to access a properties file for reading various settings.

The `vdmj.utils` package contains common utilities used by all other packages, as well as the native delegation code used by the "stdlib" VDM libraries.

2. Package Detail

2.1. vdmj

Class Summary	
VDMJ	The main class of the VDMJ parser/checker/interpreter.
VDMSL	The main class of the VDM-SL parser/checker/interpreter.
VDMPP	The main class of the VDM++ parser/checker/interpreter.
VDMRT	The main class of the VDM-RT parser/checker/interpreter.
Settings	A class holding flags for -pre, -post, -inv and -dtc, as well as the dialect.
ExitStatus	An exit status code.

Enum Summary	
Release	An enumeration for the VDM language release: classic or VDM-10.

The vdmj package contains the “main” abstract VDMJ class for the suite, plus three concrete subclasses to parse, type check and interpret a specification in the VDM-SL, VDM++ or VDM-RT dialects.

These classes just collect together a sequence of other classes to parse, type check and interpret the specification, depending on the command line arguments. The following (working) example illustrates the principles, using VDMJ classes to create a minimal interactive VDM-SL program:

```
public static void main(String[] args) throws Exception
{
    Settings.dialect = Dialect.VDM_SL;
    File file = new File(args[0]);
    LexTokenReader ltr = new LexTokenReader(file, Dialect.VDM_SL);
    ModuleReader mr = new ModuleReader(ltr);
    ModuleList modules = mr.readModules();

    if (mr.getErrorCount() == 0)
    {
        TypeChecker tc = new ModuleTypeChecker(modules);
        tc.typeCheck();

        if (TypeChecker.getErrorCount() == 0)
        {
            ModuleInterpreter interpreter =
                new ModuleInterpreter(modules);
            interpreter.init(null);
            CommandReader reader =
                new ModuleCommandReader(interpreter, "$ ");

            List<File> files = new Vector<File>();
            files.add(file);
            reader.run(files);
        }
    }
}
```

The example would be almost exactly the same for VDM++ or VDM-RT with "Class" instead of "Module" in the various names (ClassReader, readClasses, ClassTypeChecker etc.), and the dialect

constant changed from VDM_SL to VDM_PP or VDM_RT.

VDMJ has a `-o` option which causes the parsed and type-checked specification to be written out to the filename specified by the argument to `-o`. Note that the entire tree is written to the file (if there are no type check errors), so if VDMJ is invoked with four specification files, and one `-o` option, the classes specified by all four files are written to the one output file. Any input file called `"*.lib"` (rather than the more normal `*.vdmsl`, `*.vdmpp` or `*.vdmrt`) is assumed to be a pre-compiled library created by `-o` and is loaded as such without repeating the type checking. This can be faster for very large specifications. The example below shows the creation and use of `IO.lib`.

```
$ vdmpp -o IO.lib stdlib/IO.vdmpp
Parsed 1 class in 0.266 secs. No syntax errors
Type checked 1 class in 0.015 secs. No type errors
Saved 1 class to IO.lib in 0.125 secs.

$ vdmpp -i hello.vdmpp IO.lib
Loaded 1 class from IO.lib in 0.219 secs
Parsed 1 class in 0.422 secs. No syntax errors
Type checked 1 class in 0.015 secs. No type errors and 1 warning
Initialized 2 classes in 0.0 secs.
Interpreter started
> p new A().op()
Hello world!
= true
Executed in 0.015 secs.
```

2.1.1. Comments

The structure of these classes isn't very flexible. In particular, if the "load" command is given to the `CommandReader`, it is awkward to recover if parse/type errors are discovered in the new set of filenames (the `List<File>` passed to the `run` method is updated by the command reader to pass the new file names back to be parsed and checked).

The `-o` option was not as successful as I'd hoped. The idea was to serialize the tree after type checking, and load that back in quickly rather than re-parsing and re-checking the specification. Unfortunately, even though the serialization is passed through a gzip compression, the files produced are quite large and take a significant amount of time to decompress and de-serialize. For very small specifications it is certainly faster to re-parse and check them; for larger specifications there does appear to be an advantage, but not much.

Note that the serialization of the tree has to be for complete specifications (no unresolved references). This is because there is no link-editing available to combine partial specifications, and VDM++ and VDM-RT do not have a way to identify externals (VDM-SL has module imports, but VDM++ and VDM-RT have nothing similar). You can load multiple library files, or a mixture of library and source files.

2.2. vdmj.lex

Class Summary	
LatexStreamReader	A class to filter out LaTeX markup and <code>#ifdefs</code> from an input file.
BacktrackInputReader	A class to allow checkpoints and backtracking while parsing a file.
LexToken	The abstract parent class for all lexical token types.
LexBooleanToken	A class to represent a boolean token.
LexCharacterToken	A class to represent a character literal token.
LexIdentifierToken	A class to represent an identifier.
LexIntegerToken	A class to represent an integer literal token.

LexKeywordToken	A class to represent keyword tokens.
LexLocation	A class to hold the location of a token.
LexNameList	A class to hold a list of LexNameTokens.
LexNameToken	A class to hold a name.
LexQuoteToken	A class to represent a quote type token.
LexRealToken	A class to represent a real literal token.
LexStringToken	A class to represent a string literal token.
LexTokenReader	The main lexical analyser class.

Enum Summary	
Dialect	An enumeration to indicate the VDM dialect being parsed.
Token	An enumeration for the basic token types.

Exception Summary	
LexException	An exception class for lexical analyser exceptions.

The `vdmj.lex` package contains all classes concerned with the lexical analysis of specifications. This includes the lexical token reader, plus a set of value classes for representing all types of lexical token.

The base class of the lexical system is `BacktrackInputReader`, which allows a stack of markers to be pushed within a stream of characters, returning the read pointer to the previous marker when a pop operation is performed. The class allows truly random movement within the stream – which is held as an array of Unicode chars, once it has been loaded. It opens input files with a `LatexStreamReader` object (extends `InputStreamReader`), which strips out LaTeX markup and `#ifdefs`, while preserving the line numbers (turning LaTeX and `#ifdef` lines into blank lines). The `#ifdef` names available are the same as the `Dialect` constants (eg. `#ifdef VDM_PP ... #else ... #endif`). `Ifdef` statements must occur on a line by themselves.

`LexTokenStream` extends `BacktrackInputReader`. Its purpose is to make the input file look like a continuous stream of `LexTokens`. The `nextToken` method returns the next token from the stream, and `getLast` will (repeatedly) return the last token read. `Push` and `pop` mark the stream and return to a mark respectively; `unpush` removes a marker without returning to it; the `retry` method does a pop followed by a push.

The `LexLocation` class is used throughout the system to represent a location within source code, for error message reporting and code coverage. The `toString` method of the class produces a string which can be appended to another message:

```
"in <class/module> (<filename>) at <line>:<column>"
```

All value objects in the system which have a sensible position in the source code have a `LexLocation` associated with them. The class is also used to implement execution coverage tracking, with static methods to return the list of executable lines, and the locations hit or missed since the last time they were reset. It is also possible to merge coverage information from a file, written out by the `SourceFile` class (see section 2.14), which enables historical coverage to be managed.

There are several subclasses of the abstract `LexToken` to represent tokens that contain a value which is more naturally represented by a Java primitive type. For example, `LexRealToken` includes a double field, and `LexBooleanToken` contains a boolean.

The `Token` enumeration is the basic label for all token types. The enumeration includes a lookup method which, together with a `Dialect`, decides whether a given string is a token or not. Note that the dialect affects this: “static” is a token in VDM++ but is a legal variable name in VDM-SL, for example.

Lexical analysis throws a `LexException` if there are any problems. Recovery is left to the syntax analysis layer.

2.2.1. Comments

There is some modest ugliness concerned with the parsing of certain symbol sequences that look like other tokens, eg. `"mk_mod name"`. Naively, this would be a name (a `LexNameToken`) with a module part of `"mk_mod"` and the identifier part of `"name"`, but actually this is parsed as a single identifier, so that the syntax analyser can remove the `"mk_"` part and reveal the actual name.

Note that `LexLocations` have both start and end location information, though this is not used in VDMJ (it is used in the Overture editor though). The intention is to be able to identify blocks of source code that could be highlighted (eg. a whole statement or function, rather than just the start of one).

Recovering from a lexical error by raising an exception up to the syntax analysis may not be a good idea. The only recovery the syntax layer can do is to read up to some sort of safe point (eg. a semi-colon), and proceed from there. This gives comparatively poor error messages in general. It might be better to inject a likely token into the lexical stream, rather than throw an error and interrupt the stream.

`LexNameTokens` have a module/class name part and a simple name part (ie. they represent a grammatical *name* such as `C`xyz`). Whenever a name is created (as opposed to an identifier), it therefore has to include its class or module name, even if none was actually used in the specification (eg. simple parameter names are identifier patterns that are characterized by a name token, so all parameter names are held as `LexNameTokens` like `C`x`). That would be fine, except that in VDM++ and VDM-RT the presence or absence of the class qualifier in a name can have semantic significance – explicitly identifying a member in a class hierarchy for example, with `"object.X`name()"`. So `LexNameTokens` have an *explicit* flag, which means that the class name was explicitly specified by the caller, not implicitly filled in by the parser, and that flag is used during VDM++ type checking and runtime to identify the correct definition for the name. This works, but it may be over complicated. There are places where the *explicit* flag is set for very obscure reasons, which is a maintenance hazard. It may be better to take the Overture AST approach and allow names that simply don't have a class/module definition, and then take account of this when looking up definitions.

To enable VDM++ and VDM-RT function/operation overloading, `LexNameTokens` optionally include a `TypeList` qualifier, representing the parameter types of the function or operation. The qualifier, if present, is used in the `equals` method of the class, and when searching for a name in a function/operation apply - the name sought is qualified with the actual types of the arguments. The `equals` function uses the `TypeComparator` to make the test, so a function declared with an `int` parameter will match one sought with a `nat1` argument, etc. This system for managing overloading is not without its problems. Firstly, the use of the `TypeComparator` makes the `equals` method quite heavy. Secondly, it means that names cannot naively be used in Java maps because the `hashCode` of a function declaration name, and a function apply name may not be the same (if the argument types are not identical to the parameter types). Thirdly, it causes trouble when functions are applied via function variables (ie. lambda values) – such values can either be qualified with their parameter types or not, but since a function variable can either be applied (where argument types can be deduced) or just passed on (where they cannot), one or other of the name lookups will fail. To compensate, `VariableExpression` (which resolves simple names) will perform an unqualified "plain" name lookup if a qualified one fails. But a qualified lookup may succeed inappropriately by picking up an outer definition from the environment, when an inner unqualified name exists. This is tricky to solve.

2.3. vdmj.syntax

Class Summary	
SyntaxReader	The parent class of all syntax readers.
ExpressionReader	A syntax analyser to parse expressions.
TypeReader	A syntax analyser to parse type expressions.
DefinitionReader	A syntax analyser to parse definitions.
ClassReader	A syntax analyser to parse class definitions.
ModuleReader	A syntax analyser to parse modules.
PatternReader	A syntax analyser to parse pattern definitions.
BindReader	A syntax analyser to parse set and type binds.
StatementReader	A syntax analyser to parse statements.

Exception Summary	
ParserException	A syntax analyser exception.

The vdmj.syntax package contains all the syntax analysis classes. This includes eight “readers”, which implement a backtracking recursive descent parser, based on a stream of lexical tokens. The readers are all sub-classes of an abstract Reader class.

Every SyntaxReader subclass is constructed by being passed a LexTokenReader object. The reader is then responsible for returning one or more syntactic elements that it is designed to parse. For example, an ExpressionReader can be attached to a lexical stream, and be used to read an expression, or a comma separated expression list. Readers typically have methods called “read<something>” to perform the actual parse (eg. the minimal example above calls readModules from a ModuleReader).

Note that one type of reader usually needs other types of reader to complete its job. So for example, when a DefinitionReader is parsing an explicit function definition, it will use the raw lexical stream to read the function name, it will use a TypeReader to read the function’s type signature, a PatternReader to read the parameter patterns, and an ExpressionReader to read the body of the definition and any following pre or post conditions. All readers cache instances of the other readers used, and methods like getTypeReader either return the previous instance or create a new one. Note that there is no positional state information held in a reader therefore – each of them must depend on the LexTokenReader it references to (say) determine the last token read.

Several parts of the VDM grammar cannot be parsed unambiguously by reading lexical tokens in a strict sequence. For example, several parts of the grammar use a <pattern bind> symbol, which is defined to be either a pattern or a bind, but it is not possible to distinguish a pattern from a bind by looking at the next token in the stream – a type bind is a <pattern>:<type> for example, so it looks like a pattern to start with but turns out to be a bind. To overcome this, the parsers are able to backtrack. That is, the start of the <pattern bind> is marked (a push operation on the lexical stream), and an attempt is made to parse the “longest” possibility, which is a bind in this case; if that fails, the stream is popped back to the marker, and the other possibilities are tried. If all possibilities fail, there is clearly an error, though it is not clear which branch contains the error (eg. is it a pattern that is malformed or a bind that is malformed?). In this case, the parser reports the error from the branch which managed to consume the most tokens after the marker before failing. This is assumed to be the most helpful error, though it is not certain what the user intended.

The following backtrack code pattern is used frequently by the readers. Note how the ParserException is used to carry “depth” information about how far the parser progressed, and the two depths are compared at the end to see which exception to throw. The code calling this method may itself be backtracking, having pushed its own markers in the stream.

```
public PatternBind readPatternOrBind()
    throws ParserException, LexException
{
    ParserException bindError = null;

    try
    {
        reader.push();
        Bind b = readBind();
        reader.unpush();
        return new PatternBind(bind.location, b);
    }
    catch (ParserException e)
    {
        reader.pop();
        e.adjustDepth(reader.getTokensRead());
        bindError = e;
    }

    try
    {
        reader.push();
        Pattern p = getPatternReader().readPattern();
        reader.unpush();
        return new PatternBind(p.location, p);
    }
    catch (ParserException e)
    {
        reader.pop();
        e.adjustDepth(reader.getTokensRead());
        throw e.deeperThan(bindError) ? e : bindError;
    }
}
```

The top level of the parsers contain the recovery code to attempt to move the parse beyond a given syntax error. This is done by each top level case (eg. parsing a whole function definition) defining two lists of tokens: those which should be read up to, and those that should be read up to *and past* in the event of an error. Then a common recovery method ("report" in the Reader class) reads tokens up to one of those specified before continuing with the parse. The objective is, say, for a Statement reader to read tokens up to the end of the broken statement before continuing. In general this is not foolproof, and often syntax errors produce a short cascade of unrelated errors later in the specification.

2.3.1. Comments

There are some ugly parts to the parsing. One is concerned with equals definitions, which are defined as "def" <pattern bind>=<expression> "in" <expression>, but if the <pattern bind> is actually a set bind of the form "e in set S", this parses as "s in set (S = <expression>)". There is some nifty footwork in the code to get round this (see the comments in readEqualsDefinition in the DefinitionReader).

Another ugly case is concerned with object call statements, which grammatically look like object apply designators as they end in ...<name>(args). So the parser reads an apply designator, then looks inside it to see whether the object being applied is a field designator or an identifier. The former is an object member invocation, the latter is a simple operation call.

The SyntaxReader base class provides a set of methods for reading and optionally advancing by one token. The differences between them are subtle (eg. advance and return the next token, or return the current token and advance), and I suspect the code could be cleaned up by reducing the number of options.

Recursive descent parsing always has difficulty with accurate error reporting and recovery. The method chosen is not perfect.

2.4. vdmj.ast

Class Summary	
ASTConverter	The AST tree converter.

The vdmj.ast package contains one class which is used to translate an Overture AST parse tree into the equivalent structure for VDMJ. The class is constructed with a filename and an IOmlDocument obtained from the OvertureParser. A convertDocument method converts the document tree into a list of ClassDefinitions.

2.4.1. Comments

This code was originally written to import an AST from the Jflex, BYACC/J parser. Recent developments with the Overture AST will make this code obsolete, but the principles involved may be useful so the code remains for the time being. It is not used currently.

2.5. vdmj.types

Class Summary	
Type	The parent class of all static type checking types.
***Type	A *** type. There are 25 such classes.
BasicType	The parent of the basic types (numbers, booleans and characters).
NumericType	The parent of the numeric types (real, rat, int, nat, nat1)
InvariantType	A type which has an invariant function associated with it.
NamedType	A type with a name.
OptionalType	An optional type.
ParameterType	A type associated with a polymorphic parameter name.
UnionType	A union of types.
UnknownType	A type representing a parser error.
UnresolvedType	A type name identifier by the syntax analyser.
VoidType	A type indicating the absence of a type.
VoidReturnType	A type indicating that a return statement has returned "()".
PatternListTypePair	A pattern list combined with a single type.
PatternTypePair	A pattern plus a type.
TypeList	A list of types.
TypeSet	A set of types.

The vdmj.types package contains value classes that represent the various static types that can be contained in a VDM specification. They are all sub-classes of an abstract Type class.

Most simple types have a class dedicated to them of the same name, for example IntegerType or QuoteType. Similarly, the composite types have classes, like RecordType, SetType, SeqType and MapType; these have fields that in turn indicate the types of their components.

All types have a method to “resolve” themselves. The process of type resolution occurs early in the type checking process (see below), and turns UnresolvedTypes, which are just the names of types from the syntax analysis, into the actual type of the corresponding definition. The core of this happens

in the typeResolve method of UnresolvedType, though other types call typeResolve recursively for any types that they contain – for example, FunctionTypes must resolve their parameter types and return type. The type resolution mechanism contains a recursive defence to avoid types which reference themselves from blowing the stack.

All types have a method to “polymorph” themselves. This means that given an actual type parameter, they substitute that type for any ParameterTypes they contain to yield a new Type object. This is used during the type check and execution of polymorphic function instantiations, when the actual type parameters are known.

All types implement a number of “is” and “get” methods – for example, (boolean) isMap and (MapType) getMap. These are used during type checking to determine whether a type is suitable for use in (say) a map application context. For simple types, these methods return false for the “is” method, except for the type concerned, which returns “true”; and “this” from the “get” method. For more complex types, these methods support the situation where the static type checking cannot know the actual runtime type, but knows that it is one of several. In this case – the most obvious example being a UnionType – the “is” method will return true if any of the member types of the union would return true; and the “get” method will construct a new type representing the aspects of the applicable members of the union, all spliced together. The type checking can then proceed using the information of this single blend of possibilities, in the knowledge that the tests it is making could occur at runtime.

For example, if a type is a union of two records, each of which has a field called “label”, one of which is a “seq of char” and the other of which is a “nat1”, the getRecord method of the UnionType would return a new synthetic RecordType with a single field called “label”, with type “(seq of char) | nat1”. So the type checking of access to the label field would proceed as though that was its type, even though at runtime the type will be one or the other.

All types have a getAllValues method which is used during the evaluation of type binds. The method returns all the values for the type, though the only types which implements this are BooleanType, QuoteType and unions of these types; other types throw an exception if this is called.

UnknownTypes are used during error handling. Typically, an error will be encountered and reported, but rather than returning (say) the type of a sub-expression from type checking, an UnknownType is returned. This type has the property that all of its “is” methods return true, and its “get” methods will try to return a plausible Type. This means that subsequent type checking will not produce a cascade of errors as a single type checking fault deep in a specification winds its way out to the top level.

The VoidType is usually used to mean the absence of a type, so for example an operation which returns “()” would be represented by a VoidType. During the type checking of a sequence of statements in an operation, any statements which follow a “return” statement on an execution branch will be unreachable (a warning). So to distinguish this deliberate return of nothing from most statements which yield nothing, the VoidReturnType is used.

The TypeList and TypeSet classes implement lists and sets of Types, respectively. These are used in processing when a collection of types are encountered which must be turned into a single product type (TypeList) or a single union type (TypeSet). The ProductType and UnionType classes contain one TypeList and TypeSet, respectively.

2.5.1. Comments

There is no ReferenceType (compare with a ReferenceValue), yet several of the types do contain directly referenced types (like OptionalType and BracketType). It might be possible to simplify the hierarchy by adding one.

The type resolution process uses a mixture of eager and lazy evaluation. The main TypeChecker class (see below) calls the type resolution methods of the definitions in the specification, and these in turn resolve any Types that they contain. But definitions do not recurse into the expressions and statements that they contain, even though they have unresolved Types. Rather, unresolved types in expressions and statements are resolved later when the main type checking pass requires them.

Do we call them products or tuples? I went for ProductTypes and TupleValues in the end.

2.6. vdmj.expressions

Class Summary	
Expression	The parent class of all VDM expressions.
***Expression	An expression of type ***. There are >100 of these.
BinaryExpression	The parent of all binary expressions.
NumericBinaryExpression	The parent of all numeric binary expressions (+, -, *, /)
BooleanBinaryExpression	The parent of all boolean binary expressions (and, or, <=>, =>)
UnaryExpression	The parent of all unary expressions.
ExpressionList	A list of Expressions.

The vdmj.expressions package contains value classes that represent the various types of expression that can be defined in a specification. They are all sub-classes of an abstract Expression class, which defines methods for an expression to be type checked and evaluated.

The typeCheck method is implemented by all expressions, and is passed an environment defining the variables and types in scope, together with a NameScope which identifies what sorts of names are accessible (eg. whether state values are in scope – they are for expressions in operations, but not for expressions in functions).

The typeCheck method returns a Type which indicates the result of analysing the expression in the environment passed. So literal expressions simply return an appropriate type, like IntegerType. More complex expressions have to consider whether the definitions they contain affect the environment, whether those definitions have to be type checked, and whether the type check of any sub-expressions return the expected result for the overall expression.

For example, consider a “forall” expression. This will contain a bind list of variables, and a predicate to evaluate for each (at runtime). The typeCheck method of ForAllExpression is as follows:

```
@Override
public Type typeCheck(
    Environment base, TypeList qualifiers, NameScope scope)
{
    Definition def = new MultiBindListDefinition(location, bindList);
    def.typeCheck(base, scope);
    Environment local = new FlatCheckedEnvironment(def, base);

    if (!predicate.typeCheck(
        local, null, scope).isType(BooleanType.class))
    {
        predicate.report("Predicate is not boolean");
    }

    local.unusedCheck();
    return new BooleanType(location);
}
```

A MultiBindListDefinition is a type of Definition (see below) which, when given the bind list for the forall expression, can expand the Environment passed in to make the names and types of the bind variables visible. Once created, the new definition is type checked to make sure the bindings contain no errors (for example, if the bind list contains a set bind, the expression representing the set must be a SetType).

A new FlatCheckedEnvironment is created to chain the new definitions onto the base Environment passed in, and this is used to type check the predicate of the forall expression. The return value of this is the Type of the predicate, which must be a boolean expression. The local environment is then checked to see whether all the names added to it by the bind list were actually used when type

checking the predicate; any unused variables generate a warning. Lastly, this method returns a boolean Type, since a forall expression returns a boolean.

The typeCheck method on Expression is also passed a TypeList. This is used when trying to resolve name overloading during function and operation application in VDM++ and VDM-RT. The typeCheck method of ApplyExpression starts by generating a TypeList from the typeCheck of each of the argument expressions it has. That may generate (say) [int, int, bool]. That list is then passed to the typeCheck method for the root of the apply (the thing being applied). If this root is a VariableExpression or a FieldExpression, the name of the variable or field is qualified with the list of types passed in, and this is used to find an overloaded name of a function or operation definition that has parameters whose types are compatible with the arguments (note, compatible with, not identical to). If it turns out that the variable or field is actually a map (which has no qualifiers), the search is repeated for the name without type qualification.

The other important method on Expressions is the eval method. This is called to evaluate the expression given the runtime Context (the runtime equivalent of an Environment). The method returns a Value object, which can represent any value in VDM. The eval method for PlusExpression is as follows:

```
@Override
public Value eval(Context ctxt)
{
    breakpoint.check(location, ctxt);

    try
    {
        double lv = left.eval(ctxt).realValue(ctxt);
        double rv = right.eval(ctxt).realValue(ctxt);

        return NumericValue.valueOf(lv + rv, ctxt);
    }
    catch (ValueException e)
    {
        return abort(e);
    }
}
```

All Expressions and Statements contain a Breakpoint object, and the eval method of all expressions and statements call their breakpoint's check method at the start. This usually does nothing, unless a breakpoint is set at this location, in which case execution stops and calls debugger code.

The PlusExpression evaluates its left and right hand sides, and converts the resulting Value objects to raw Java doubles. The result is a new Value, created using the valueOf method of NumericValue, which will create the simplest type of NumericValue capable of holding the result of the addition – for example, this might be an IntegerValue (for -123) or a NaturalOneValue (for +123) or a RealValue (for 1.23).

Note that the code may throw a ValueException, and that this is caught within the eval method rather than being propagated. This can only occur in the conversion of the sub-expression results to doubles, or the construction of the result Value. ValueExceptions indicate problems while evaluating an expression, but they are caught and propagated using the abort method, which creates and throws a ContextException (a Java RuntimeException). This is done to distinguish between expected value errors – for example, trying to convert a Value to one of several types in a union, and failing before the right one is found – and serious errors, which should cause the system to halt. The value system does not have location information (what is the location of the pure value "hello"?), and so this double-propagation of the exception also allows the value error to be located in an expression that caused it.

The most complex evaluations are for functions and operations, but these are delegated to the corresponding FunctionValue and OperationValue classes. This is so that a function value can be separately created (for example via a lambda expression) and applied. Operations cannot be created like this, but having operation values too means that, in VDM++ or VDM-RT, an object becomes a map of names to values – whether those are instance variable values or member operations and functions. The eval method of the ApplyExpression to call a function is therefore just:

```
try
{
    Value object = root.eval(ctxt).deref();

    if (object instanceof FunctionValue)
    {
        ValueList argvals = new ValueList();

        for (Expression arg: args)
        {
            argvals.add(arg.eval(ctxt));
        }

        FunctionValue fv = object.functionValue(ctxt);
        return fv.eval(argvals, ctxt);
    }
    else if (object instanceof OperationValue)
    ...
}
```

All Expressions implement a method called `findExpression`. This has a line number parameter, and all implementations are responsible for returning themselves if they start on the line number, or recursing into their sub-expressions if they have them. This is used to set breakpoints on specific lines of a specification.

2.6.1. Comments

The implementation of name overloading may be problematic. A `LexNameToken` is optionally qualified with a `TypeList`, and the `equals` method uses the `TypeComparator` (part of `vdmj.typechecker`) to make a compatible comparison of the two names' type lists. Care must be taken when comparing names, especially during the "bootstrap" phases when qualifiers are not yet available.

2.7. vdmj.statements

Class Summary	
Statement	The parent class of all statements.
***Statement	A statement of type ***. There are 40 or so of these.

The `vdmj.statements` package contains value classes that represent the various types of statement that can be defined in a specification. They are all sub-classes of an abstract `Statement` class, which defines methods for a statement to be type checked and executed.

Many of the principles for Statements are the same as those for Expressions covered above. Like Expressions, all Statements include a `typeCheck` method which is passed an `Environment` and a `NameScope` – though note that there is no need for a `TypeList` of qualifiers because an operation call can only be rooted on something that is already known by the statement (an operation name or an object designator), whereas function application in an expression has to evaluate an arbitrary expression to generate the root to which to apply the arguments.

Similarly, like Expressions, all Statements define an `eval` method which executes them, returning a `Value` – though statement executions usually return `VoidValues`.

Statements implement a method called `exitCheck`, which explores the statement tree (following blocks and branches from compound statements) looking for statements which can raise an exit status, like the `ExitStatement` itself. This is used in the type checking of statements that catch and process exits, like `TrapStatements`. The method returns a set of exit Types that can be thrown.

All Statements implement a method called `findStatement`. This has a line number parameter, and all

implementations are responsible for returning themselves if they start on the line number, or recursing into their sub-statements if they have them. This is used to set breakpoints on specific lines of a specification.

2.7.1. Comments

The exitCheck is not perfect. In particular, the check does not cross the boundary of an operation call from a statement block, nor can it follow an expression evaluation that involves operation calls. That would require code to work out which operation(s) are actually involved, and it would require recursive defence against operations which recurse. So exitCheck produces false negatives (indicates that operations can't exit, when in fact they can). I gather VDMTools is better, but not perfect either. It is known to produce false positives.

2.8. vdmj.patterns

Class Summary	
Pattern	The parent type of all patterns.
****Pattern	A pattern of type ****. There are 14 such pattern types.
Bind	The parent class of SetBind and TypeBind.
MultipleBind	The parent class of MultipleSetBind and MultipleTypeBind.
PatternBind	A pattern or a bind.
PatternList	A list of patterns.

The vdmj.patterns package includes value classes to represent the various patterns and binds in a specification. They are all subclasses of an abstract Pattern, Bind or MultipleBind class.

Patterns generate definitions, given a Type. For example, the pattern “[a,b,c]” will produce definitions for the three integer variables, given that it is of type “seq of int”. The process of definition generation is recursive over the tree of nested patterns that might be defined. So the getDefinitions(Type) method which all patterns implement, will recurse for those pattern types which are defined as containing sub-patterns. The leaves of the pattern tree are the simple pattern types: BooleanPattern, CharacterPattern, etc. At the leaves, it is only IdentifierPattern which produces variable definitions.

Similarly, patterns generate a list of name/value pairs given a Value. The value is matched against the “shape” of the pattern and its sub-patterns, and any IdentifierPatterns at the leaves are populated with the corresponding part of the original value. This getNamedValues method is called from definitions which include patterns (such as function parameter definitions) when the actual values are required.

Patterns can also be asked for a simple list of their variable names, which involves a depth search for IdentifierPatterns.

Patterns include a typeResolve method because ExpressionPatterns can contain references to UnresolvedTypes that need to be resolved early in the type check. Definitions which include Patterns recurse into the pattern's typeResolve from their own typeResolve methods.

A Bind, which is sub-classed by SetBind and TypeBind, comprises a Pattern and a set of Values (either an explicit set, or the set of all the Values that a Type can generate, in theory). These are used in quantified expressions, like “exists {a,b} in set S & a.type = b.type”. Here the set pattern {a,b} contains two identifier patterns that must (potentially) iterate through all the elements of S, taking the corresponding values. To drive the iteration, the Bind needs to generate all the possible Values, which are then given to the pattern to generate name/value pairs for each iteration. Therefore Bind has a method called getAllValues, which is implemented by both sub-classes (the TypeBind implementation calls the getAllValues method of the Type, which is an error for everything except BooleanType, QuoteType and unions of these).

MultipleBind, which is sub-classed by MultipleSetBind and MultipleTypeBind, is very similar except that

they comprise a list of patterns and a set or type. But they still have a `getAllValues` method which collects together all the possible values in the set.

Note that in VDMJ, the generation of all values from a set includes the permutations of all the orderings of the values in that set. Internally, sets of Values are held in a `ValueSet` which is actually an ordered list (but with set semantics, with regard to no duplicates). Therefore the `getAllValues` methods for the various set binds call the `permuteSets` method of `ValueSet` (via the same method on `SetValue`, which preserves the ordering). Note also that the original set is sorted before this process, which means that given the same set content, the order of processing in a bind (and therefore any looseness based on it) is deterministic.

```
@Override
public ValueList getBindValues(Context ctxt)
{
    try
    {
        ValueList results = new ValueList();
        ValueSet elements = set.eval(ctxt).setValue(ctxt);
        elements.sort();

        for (Value e: elements)
        {
            e = e.deref();

            if (e instanceof SetValue)
            {
                SetValue sv = (SetValue)e;
                results.addAll(sv.permutedSets());
            }
            else
            {
                results.add(e);
            }
        }

        return results;
    }
    catch (ValueException ex)
    {
        abort(ex.getMessage(), ctxt);
        return null;
    }
}
```

2.8.1. Comments

Binds and `MultipleBinds` look like they have a lot in common and could probably be put together to avoid a small amount of code duplication.

2.9. vdmj.traces

Class Summary	
TraceDefinition	An abstract class representing a trace definition.
TraceDefinitionTerm	A class representing a sequence of trace definitions.
TraceLetBeStBinding	A class representing a let-be-st trace binding.
TraceLetDefBinding	A class representing a let-definition trace binding.
TraceRepeatDefinition	A class representing a repeated trace definition.

TraceCoreDefinition	Abstract of all core trace expressions.
TraceApplyExpression	A class representing a core trace apply expression.
TraceBracketedExpression	A class representing a core trace bracketed expression.
TraceNode	An abstract class representing an expansion node.
AlternativeTraceNode	An expansion node for alternatives.
RepeatTraceNode	An expansion node for repeats.
SequenceTraceNode	An expansion node for sequences.
StatementTraceNode	An expansion node (leaf) for statement applies.
TestSequence	A sequence of CallSequences
CallSequence	A sequence of Statements.
Permutor	A utility to permute a set of values.
TraceVariable	A class containing a name/value/location tuple.
TraceVariableList	A list of TraceVariables
TraceVariableStatement	A statement wrapping a TraceVariable.

Enum Summary	
Verdict	A test outcome: PASSED, FAILED or INDETERMINATE.
TraceReductionType	A reduction type, RANDOM or various SHAPE types

The vdm.traces package includes classes to represent the definitions which can occur in a VDM-SL, VDM++ or VDM-RT "traces" section, and their subsequent expansion and execution.

The subclasses of TraceDefinition are uncontroversial and follow the structure of the trace grammar closely. These classes have the usual typeCheck methods which permit the trace specifications to be checked as part of the overall specification type check phase.

In order to evaluate traces, they must first be expanded into all the possible execution paths represented by the definition. For example, a trace that is of the form "a;(b|c);d" would expand to "a;b;d" and "a;c;d". Similarly, "a{1,3}" would expand to "a", "a;a" and "a;a;a". Repeats, sequences and alternations multiply together to generate large numbers of tests very quickly – this is the whole point of combinatorial test specification. The TraceNode class and its subclasses represent the expanded traces, and are generated by "expand" methods on all TraceDefinitions. Strictly, these classes don't expand the traces, but produce a tree structure that is capable of expanding them. The getTests method of TreeNodes actually expands the tests, returning a TestSequence, which is a list of CallSequence (ie. a list of tests), and a CallSequence is a list of Statements (CallStatements, CallObjectStatements or TraceVariableStatements).

All TraceNodes include a field called variables, which may contain a TraceVariableList. These are populated by TraceLetBeStBinding and TraceLetDefBinding when a particular name/value pair is chosen in the expansion of a give test sequence. The getVariables method of TraceNode returns a CallSequence which is populated with any TraceVariableStatements for the expansion. When executed, these statements just add their named value to the local context, making it available to subsequent calls in the test.

A top level NamedTraceDefinition (see below) is created for each parsed trace definition, including a method to extract all the tests from the definition tree it contains.

The execution of a trace is coordinated from a method called runtrace on the Interpreter object (see below). This first gets a TestSequence from the NamedTraceDefinition. Then for each CallSequence in the TestSequence, it initializes the interpreter, and calls its runOneTrace method, passing the CallSequence. The runOneTrace method executes the CallSequence passed and returns a list of java.lang.Object, being either the return Values from the test steps or error messages, where the last item will always be an instance of a Verdict object, indicating the test outcome.

Tests which have a FAILED verdict are "stemmed" by runtrace such that remaining tests in the TestSequence which have the same stem are marked as "filtered" by this test – ie. there is no point in running them because their initial sequence of calls will fail at the same point, for the same reason. Before each test is executed, its filtered flag is tested, and such tests are not executed. Note that the stem check includes the value of TraceVariableStatements in the test. This means that tests which are superficially the same, but which have different variable values will not match.

The runtrace and runOneTrace methods have a boolean debug argument. If true, this allows errors in trace execution to trap into the debugger, otherwise such errors are caught and returned in the list of results without interrupting the execution of the trace.

The TestSequence class has a method called reduce, which can be used to cut down a very large collection of tests before execution. The method is passed three parameters via runtrace: the proportion of the tests to keep (a float value between 0 and 1), a TraceReductionType which indicates the type of reduction to perform, and a random seed value. The simplest reduction type is RANDOM, which seeds a PRNG and then randomly removes tests until the required number remain. The weakness of this type of reduction is that it does not guarantee to preserve all the possible "shapes" of operation call sequences in the test collection, a shape being a sequence of named operation calls (regardless of the arguments passed). The SHAPES reduction type guarantees to keep at least one test of each test shape. There are three subtypes, SHAPES_NOVARS which ignores TraceVariableStatements when determining shapes, SHAPES_VARVALUES which takes variable names into account but not their values, and SHAPES_VARS which considers variables and their values.

2.9.1. Comments

There is a minor quibble in the parsing of trace sections, in that the grammar prohibits the use of semi-colons between trace definitions (while permitting them to separate parts of a trace definition). VDMJ permits these separate semi-colons; the Overture parser currently does not.

I tried very hard to combine the expansion of the tests with the classes that represent the parsed trace definitions – this is just a tree after all. But I couldn't get it working properly, hence the solution where the definitions are expanded into a separate tree, which is then "walked" to generate the trace permutations.

The intention of the TraceVariable was to hold information about the location of a particular value from a set of (possibly) anonymous values, such as "let x in set {new A(1), new A(2), ...} in ..." Here, x will take one of the values from the set in each expansion, but it will always be called "x" and it is hard to distinguish cases when a given test fails. Unfortunately, this is very hard to achieve without giving (pure) Values a location. Currently, the location stored with TraceVariables is the location of the name of the variable. The name/value is included in the CallSequence (via the TraceVariableStatement), so debuggers can look at the raw value, which may be of some help.

2.10. vdmj.definitions

Class Summary	
Definition	The abstract parent of all definitions.
AccessSpecifier	A class to represent a [static] public/private/protected specifier.
AssignmentDefinition	A class to represent assignable variable definitions.
ClassDefinition	A class to represent a VDM++ or VDM-RT class definition.
SystemDefinition	A class to represent a VDM-RT system class definition.
CPUClassDefinition	A class to represent a VDM-RT CPU definition.
BUSClassDefinition	A class to represent a VDM-RT BUS definition.
ClassInvariantDefinition	A class to hold a class invariant definition.
EqualsDefinition	A class to hold an equals definition.

ExplicitFunctionDefinition	A class to hold an explicit function definition.
ExplicitOperationDefinition	A class to hold an explicit operation definition.
ExternalDefinition	A class to hold an external state definition.
ImplicitFunctionDefinition	A class to hold an implicit function definition.
ImplicitOperationDefinition	A class to hold an explicit operation definition.
ImportedDefinition	A class to hold an imported definition.
RenamedDefinition	A class to hold a renamed import definition.
InheritedDefinition	A class to hold an inherited definition in VDM++.
InstanceVariableDefinition	A class to hold an instance variable definition.
LocalDefinition	A class to hold a local variable definition.
MultiBindListDefinition	A class to hold a multiple bind list definition.
MutexSyncDefinition	A class to hold a mutex synchronization definition.
PerSyncDefinition	A class to hold a permission synchronization definition.
StateDefinition	A class to hold a module's state definition.
ThreadDefinition	A class to hold a thread definition.
TypeDefinition	A class to hold a type definition.
UntypedDefinition	A class to hold a definition of, as yet, an unknown type.
ValueDefinition	A class to hold a value definition.
NamedTraceDefinition	A class to hold a named trace definition.
ClassList	A class for holding a list of ClassDefinitions.
DefinitionList	A class to hold a list of Definitions.
DefinitionSet	A class to hold a set of Definitions with unique names.

The `vdmj.definitions` package contains classes representing the definitions in a specification. They are all sub-classes of an abstract `Definition` class.

All definitions have a few things in common (fields of the abstract class). They belong to a `Pass`, which guides the type checking; they have a `location`; they have a `name`, though this may be null if they contain sub-definitions; and they have a `NameScope` to define what sort of name(s) they define, which compliments the name scope used in type checking that searches for names of certain types.

Definitions define a `typeResolve` method which is used very early on to resolve the `UnresolvedTypes` that may have come through from the syntax analysis. For example, an `ExplicitFunctionDefinition` would `typeResolve` the `Type` of the function, and if there were pre or postconditions, these expressions would be `typeResolved`, and any parameter patterns would be `typeResolved`.

Definitions also define a `typeCheck` method, which is similar to the ones defined for `Expression` and `Statement`, except that there is no `Type` to return. For example, the `ExplicitFunctionDefinition` performs the following tasks in its `typeCheck` method:

- If there are any polymorphic type parameters for this function, check that the overall function type does not reference any type parameters except those named type parameters.
- For each type parameter, create a `LocalDefinition` of a `ParameterType` and add this to a local `Environment`.
- Check that the parameter patterns match the overall `Type`'s parameters, and iterate through curried sets of parameters, using the return value from the overall `Type` (and its return value and so on for subsequent sets of parameters). Remember the expected result.

-
- Extend the local Environment with definitions for all the variables of all the patterns from all of the curried parameter sets.
 - Type check the definitions this produced in the base environment (this will just do type resolution, if necessary).
 - Label the local Environment as static (VDM++) if the definition's access specifier is static.
 - If we are in VDM++ or VDM-RT and the function is not static, add a "self" definition to the local Environment.
 - If there is a precondition expression, type check the definition for it.
 - If there is a post condition expression, type check the definition for that too.
 - Type check the body expression of the function, remembering the actual type returned.
 - If the actual return type is not assignable to the expected return type, raise an error.
 - If the VDM++/VDM-RT accessibility of the expected return type is narrower than that of the definition itself, raise an error (eg. a public function cannot have a private return type).
 - If the function is recursive and does not define a "measure" function, raise a warning, else if there is a measure defined, check that it exists and has the correct type.
 - Check that the parameter variables have been referenced in the local Environment, else raise an unused parameter warning. (This is suppressed for pre and post conditions, which are permitted to not necessarily use their implicit parameters).
 - Return.

This illustrates the principles that are used by all definitions' typeCheck methods.

Some definition types are only used to "wrap" others. For example, during module imports and exports, a definition may be imported and/or renamed. These methods just delegate their calls to the referenced definition that they wrap.

As with Patterns, definitions can yield their contained definitions or a list of names or name/value pairs that they define. This is the purpose of the getDefinitions, getVariableNames and getNameValuePairs methods. Simple local definitions only define one variable, but many definition types include a Pattern specifier that may define many variables.

The findName method is implemented by all definitions to return whether they define a name being sought by type checking. As above, for simple definitions, this just compares their name and scope with that being searched for, but for definitions that include patterns, all the names generated by the pattern must be considered.

The findType method is implemented by those definitions that define a type (TypeDefinition, StateDefinition and ClassDefinition).

Definitions also define findExpression and findStatement methods which recurse into their bodies in search of an expression or statement that starts on the given line.

The ClassDefinition class is slightly different from the others in that its main job is to contain the definitions in a class, though it does have the job of setting up the static and instance environment for new objects when a "new Object()" statement is executed. It is also responsible for stitching together the class hierarchy and arranging for symbols to be inherited so that type checking may be performed. The hierarchy is built during the generation of implicit definitions.

Note that class static data (eg. instance variables that are declared static) is held inside the ClassDefinition at runtime (in the public/privateStaticValues fields), whereas object instance data is held inside an ObjectValue (produced by the makeInstance method of ClassDefinition), though references to the static data is included in the object's member list, so that the runtime can find them.

See section 2.14 for more information about runtime variable access.

The `SystemDefinition` class is a subclass of `ClassDefinition`, and adds VDM-RT specific processing for system classes. The implicit definition generation (see section 2.12) checks whether the definitions in the system class meet the VDM-RT restrictions. A `systemInit` method is called during system initialization and creates the necessary CPU and BUS objects from the system definition. The `newInstance` method is overridden, since it is not legal to create an instance of a VDM-RT system class.

The `CPUClassDefinition` and `BUSClassDefinition` classes represent VDM-RT CPU and BUS classes respectively. These are also subclasses of `ClassDefinition`, and override the `newInstance` method to perform special processing. Both classes create their operations (ie. create `ExplicitOperationDefinitions` for their definition list) by parsing a string literal representing the operations required. For example:

```
private static String defs =
    "operations " +
    "public BUS:(<FCFS>|<CSMACD>) * real * set of CPU ==> BUS " +
    "    BUS(policy, speed, cpus) == is not yet specified;";

private static DefinitionList operationDefs()
    throws ParseException, LexException
{
    LexTokenReader ltr = new LexTokenReader(defs, Dialect.VDM_PP);
    DefinitionReader dr = new DefinitionReader(ltr);
    dr.setCurrentModule("BUS");
    return dr.readDefinitions();
}
```

This technique permits the rest of the code to use these operations as normal. The *is not yet specified* processing intercepts the operation calls for CPU and BUS, calling back to the `CPUClassDefinition` and `BUSClassDefinition` classes to perform the actual processing.

2.10.1. Comments

There is a great deal of commonality between some definitions, especially implicit and explicit functions and operations. It might be possible to simplify the code by creating abstract bases for these.

2.11. vdmj.modules

Class Summary	
Export	The parent class of all export declarations.
Export***	A class for representing exports of a given type.
Import	The parent class of all import declarations.
Import***	A class for representing imports of a given type.
Module	A class holding all the details for one module.
ModuleList	A list of Modules.

The `vdmj.modules` package contains value classes that represent the modular structure of VDM-SL specifications, including their import and export declarations.

The only purpose of these classes is to represent the parsed module structures from the specification, and to generate/find the list of exported and imported definitions that extend the scope of what is visible from a single module.

The ModuleList class is important because it contains the "initialize" method which is used to set the initial state of all modules when the interpreter is started.

2.12. vdmj.typechecker

Class Summary	
TypeChecker	The abstract root of all type checker classes.
ClassTypeChecker	A class to coordinate all class type checking processing.
ModuleTypeChecker	A class to coordinate all module type checking processing.
Environment	The parent class of all type checking environments.
FlatEnvironment	Define the type checking environment for a list of local definitions.
FlatCheckedEnvironment	Define the type checking environment for a list of local definitions, including a check for duplicates and name hiding.
ModuleEnvironment	Define the type checking environment for a modular specification.
PrivateClassEnvironment	Define the type checking environment for a class as observed from inside.
PublicClassEnvironment	Define the type checking environment for a set of classes, as observed from the outside.
TypeComparator	A class for static type checking comparisons.

Enum Summary	
NameScope	An enum to represent name scoping.
Pass	An enum to indicate which type checking pass a definition belongs to.

The vdmj.typechecker package includes classes which organize the type checking of VDM-SL, VDM-RT and VDM++ specifications. Most of the actual type checking is performed by the Definition, Expression and Statement subclasses that collectively describe the specification (above), but the typechecker package defines the supporting classes to invoke the type checking methods of the other objects, and to represent the static environment in which type checking is performed.

Type checking is different for the three dialects, though the checking of the basic statements and expressions is very similar, and VDM-RT is really an extension of VDM++. Therefore there is one common abstract TypeChecker class with two subclasses: ModuleTypeChecker and ClassTypeChecker. The subclasses are constructed with a list of modules or classes – which is the overall result of a successful syntax analysis – and they implement a single abstract method from their parent, called typeCheck. The method takes no arguments and returns no result. Any errors or warnings raised during type checking are recorded by the parent class (see VDMMessage).

The sequence of events is slightly different for the type check of modules and classes, but they follow the same principles. For modules, the sequence is:

- Check for duplicate module names in the list passed
- For each module, generate its definitions' implicit definitions (like pre and post functions)
- For each module, check the export definitions exist and are of the declared type, and make a list of exported definitions for the module.
- For each module, go through the import definitions and resolve against the exports.

- Create a list of all definitions from all modules (including their imports), create an Environment that contains them all, and attempt to perform type resolution on them – ie. find the type definition for every named type.
- In the pass order: [types, values, definitions], for each module, create a ModuleEnvironment representing the visible definitions, and type check the definitions of the given pass. This calls the typeCheck method on the definitions, which calls the similar method on the definitions subparts, if any.
- Report any discrepancies between the final checked types of the modules' definitions and their explicit imported types.
- Any definition names that have not been referenced or exported produce "unused" warnings.

There are a couple of important points to note:

Firstly, the syntax analysis does not understand anything about the relationship of a type name in a declaration to its type definition. All type names come through from the syntax phase as UnresolvedTypes, which simply have a name. So a very early phase of the type checking must find the corresponding type definition, in order to understand the structure of the type and what is/is not a legal type manipulation. This is done on a global basis, even though not all types are in scope for a module (all type names are fully qualified with a module name, so there is no ambiguity). A subsequent pass, which uses just those definitions that are visible, will subsequently spot any scope problems.

Secondly, "environments" are used to support the type checking. An environment (a subclass of the abstract Environment class) is essentially a list of names and corresponding definitions that are in scope at any point. The different subclasses allow the different scope rules to be followed, so for example a module's type checking will start with a ModuleEnvironment that references a single module definition, and understands the rule about resolving names from its imported definitions and its own definitions. Different environments are then chained together, so when the module environment is passed to (say) a function definition, the type checking creates a new environment with local definitions for the names that are generated by the parameter patterns. Since the parameter names are in scope for the body of the function, the chain of two environments is passed to the type checking of the body expression. That may in turn involve "let" expressions that define further local variables that are chained onto the environment before their bodies are type checked, and so on. As the chain unwinds (as typeCheck methods return), each environment in the chain is checked to see whether all of its definitions were referenced; any that were not referenced generate "unused" warnings.

The two most important methods on an Environment subclass are findName and findType, which lookup definitions by name (using the same methods on the Definition classes they reference). There is also a name scope parameter passed to findName to indicate what sorts of names are sought – for example, functions "see" value and parameter names, operations see these names and state variables, and the post conditions of operations see names, state and "old" names. The name scope (mask) is passed around the tree of typeCheck invocations as the names that are visible in a given context is generally only known by the caller – eg. an expression does not know that it is part of an operation, so it must be told from the outside that state variables are in scope.

Very similar principles are followed by the ClassTypeChecker, which performs the following actions:

- Make sure there are no duplicate class definitions.
- For all classes and their definitions, generate the implicit definitions. This includes the construction of the class type hierarchy and the implicit local names for access to inherited definitions. VDM-RT specifications limit what can be done in system classes here.
- Create a public class environment that can see all public class definitions.
- For each class, chain a private class environment to the public environment, and perform type resolution on the definitions in the class.

- For each class, check for overloading and overriding of its definitions.
- In the pass order: [types, values, definitions], for each class, create a private class environment, and type check the definitions of the given pass.
- Check for any definition names that have not been referenced or exported, and produce "unused" warnings.

The use of public and private class environments to control the resolution of names is exactly analogous to the module case, though the class versions use the static/public/protected/private definition modifiers to decide on visibility.

The TypeComparator is used at the heart of type checking. The "compatible" method is used to decide whether two Types are assignment compatible (with "possible" semantics). This involves finding the "underlying" type (eg. removing the names of types) and making flexible recursive comparisons where unions are involved. It also has to have recursive defence, since type structures may reference themselves. Two optional types are always considered compatible (as they could both be nil). Compound types that involve more than one subtype (sets, sequences, maps, functions and operations, records and classes) are unpicked and their subtypes recursively compared. Finally, a Type.equals comparison is made between simple types.

The TypeComparator is also involved in proof obligation generation for subtype testing (see below. This is effectively a "definite semantics" check, since PO generation is concerned with covering the gaps that "possible semantics" leaves). The "isSubType" method takes two types and returns a boolean indicating whether the first is a subtype of the second – for example, a nat1 is a subtype of a real (all nat1 values are real values), but not the other way round. With union types, a simple type is a subtype of a union if it is a subtype of any of the union members; a union type is a subtype of another union if every member of the first union is a subtype of the second union.

2.12.1. Comments

I had a lot of trouble getting the order of the initialisation and type checking right to be able to deal with all the specifications in the test suite. I'm not 100% sure that there aren't still obscure orderings of declarations that will defeat it.

The TypeComparator is currently a static class, which means that its methods need to be synchronized to work with VDM++/VDM-RT threads (it has state for recursion defence). This isn't really necessary, but there are quite a few places where the code would have to create a new TypeComparator instance if this is changed.

Type checking error messages are produced by static methods on the TypeChecker class, and the error count is held statically there too. This is because it is otherwise difficult to find the type checking object instance deep in a typeCheck call chain. This is in contrast to the syntax analysers, where a Reader keeps track of the errors for itself (plus any from the Readers it creates).

2.13. vdmj.pog

Class Summary	
ProofObligation	The abstract root of all proof obligations.
***Obligation	A particular type of proof obligation.
ProofObligationList	A list of proof obligations.
POContext	The abstract root of all obligation contexts.
PO***Context	A particular type of obligation context.
POContextStack	A stack of obligation contexts.

Enum Summary	
POType	An enumeration of the various proof obligation types.

The vdmj.pog package defines classes that support the generation of proof obligations. Most of the actual obligation generation is performed by Definitions, Expressions and Statements, which include a getProofObligations method that returns a ProofObligationList.

A typical proof obligation contains a stack of nested "contexts" in which a particular obligation must be determined, plus a specific obligation to check. Therefore there are two distinct class hierarchies in the pog package: the POContext hierarchy, and the ProofObligation hierarchy.

For example, if "m" is a map of int to int, a simple function like:

```
f: int -> int
  f(i) == if i < 10 then m(i) + 1 else m(i) - 1;
```

would generate two proof obligations, requiring that the two m(i) map applications are correct in the "then" and "else" branches, respectively:

```
A`f(int): map apply obligation in 'A' (test.vpp) at line 15:32
(forall i:int &
  ((i < 10) =>
    i in set dom m))
```

```
A`f(int): map apply obligation in 'A' (test.vpp) at line 15:46
(forall i:int &
  (not (i < 10) =>
    i in set dom m))
```

Notice that both obligations contain an outermost "forall" that represents the possible function parameter values, and that the "if" test value is determined to be true or false in order to check the map application in the two branches. These two form the "context" of the proof obligation; the last line in both POs is the actual obligation, which tests that the argument is within the domain of the map.

The outer forall context is represented by a POFunctionDefinitionContext object, and the if/else contexts are represented by POImpliesContext and PONotImpliesContext objects, respectively. The proof obligation itself is a MapApplyObligation. The two contexts form a stack for each obligation, and these are held by a POContextStack, which extends Stack<POContext>.

To generate proof obligations for an entire specification, an empty POContextStack is created and the getProofObligations method of each definition is called, passing the stack. Depending on the definition type, the implementation may add to the context (eg. a function definition would push a POFunctionDefinitionContext) before calling getProofObligations for their inner expression(s) or statement(s). In the example above, the getProofObligation method of the IfExpression that comprises the function body would be called. That in turn would push an POImpliesContext before generating obligations in its "then" sub-expression, **popping** the stack, and pushing a PONotImpliesContext before generating obligations for everything in the else-if list, and final else. Lastly it would pop the stack back to the state it found it in before returning a list of all the generated POs.

```
public ProofObligationList getProofObligations(POContextStack ctxt)
{
    ProofObligationList obligations = ifExp.getProofObligations(ctxt);

    ctxt.push(new POImpliesContext(ifExp));
    obligations.addAll(thenExp.getProofObligations(ctxt));
    ctxt.pop();

    ctxt.push(new PONotImpliesContext(ifExp));           // not (ifExp) =>
```

```

        for (ElseIfExpression exp: elseList)
        {
            obligations.addAll(exp.getProofObligations(ctxt));
            ctxt.push(new PNotImpliesContext(exp.elseIfExp));
        }

        obligations.addAll(elseExp.getProofObligations(ctxt));

        for (int i=0; i<elseList.size(); i++)
        {
            ctxt.pop();
        }

        ctxt.pop();

        return obligations;
    }

```

Notice that the `IfExpression` does not actually generate any proof obligations itself; it only sets up context so that obligations generated from its sub-expressions will be correct. The `ApplyExpression` actually generates the proof obligations in this example:

```

public ProofObligationList getProofObligations(POContextStack ctxt)
{
    ProofObligationList obligations = new ProofObligationList();

    if (type.isMap())
    {
        MapType m = type.getMap();
        obligations.add(
            new MapApplyObligation(root, args.get(0), ctxt));

        Type atype = argtypes.get(0);

        if (!TypeComparator.isSubType(atype, m.from))
        {
            obligations.add(new SubTypeObligation(
                args.get(0), m.from, atype, ctxt));
        }
    }
    ...
}

```

Here, the apply expression is first tested for whether it is a map application (it could be a function or operation application), and if so, a new `MapApplyObligation` is created, which is passed the context. Similarly, if the apply argument type is not a subtype of the domain of the map, a `SubTypeObligation` is also generated.

The constructor of the `ProofObligations` generated use the context passed to generate the "value" of the obligation (ie. the string form of its expression):

```

public MapApplyObligation(
    Expression root, Expression arg, POContextStack ctxt)
{
    super(root.location, PType.MAP_APPLY, ctxt);
    value = ctxt.getObligation(arg + " in set dom " + root);
}

```

The `getObligation` method on the context stack will generate a string composed of the contexts in the stack, plus the string passed in for the obligation. It is also responsible for the indentation and bracketing of the expressions. All `ProofObligation` subclasses are similar to the example above, though the more complex ones take a lot of effort to generate the string of the obligation. The most complex obligation is `SubTypeObligation`, which has a recursive private method to generate all the subtype tests that are required for the "structure" of the type being considered.

The process of proof obligation generation is exactly analogous with operation definitions which include statements – though often, statement obligations are generated without any context since in general it is too difficult to determine the scope of side effects generated by a specification.

2.13.1. Comments

It is possible that there are other obligation types that should be added, especially in the case of VDM++/VDM-RT.

ProofObligation subclasses generate the entire string of the obligation, including the context in which it is generated. This seemed better than keeping the context objects with the obligation, but it does mean that the structure of the PO is lost in the flat string.

The generation of expression strings depends on the accuracy of the toString methods for Expressions. These have been tidied up, but unfortunately preserving the precedence of the original operators means that many expression strings end up being excessively bracketed.

2.14. vdmj.runtime

Class Summary	
Interpreter	An abstract VDM interpreter.
ModuleInterpreter	The VDM-SL module interpreter.
ClassInterpreter	The VDM++ and VDM-RT interpreter.
Context	A class to hold runtime name/value context information.
RootContext	An abstract context for the root of a function or operation call.
ObjectContext	A root context for object member invocations.
StateContext	A root context for non-object member invocations.
ClassContext	A root context for static member invocations.
Breakpoint	The concrete root of the breakpoint hierarchy.
Stoppoint	A breakpoint where execution must stop.
Tracepoint	A breakpoint where something is displayed, but execution continues.
SourceFile	A class to hold a source file for source debug output.
ThreadState	A class to hold runtime information for a VDM thread.

Exception Summary	
ContextException	A fatal interpreter error, including a “stack” context to dump.
DebuggerException	An exception used to stop the interpreter cleanly from the debugger.
ExitException	An exception used to implement exit statements.
PatternMatchException	A non-fatal exception indicating a pattern match has failed.
ValueException	A non-fatal exception concerning an evaluation.

The vdmj.runtime package defines the abstract interpreter, and its subclasses to interpret VDM-SL, VDM++ and VDM-RT specifications. It includes classes to represent the runtime execution context, exceptions which can be generated at runtime, and classes for breakpoint handling.

At its simplest, an interpreter has to create a runtime environment that represents the globally visible name/values in a specification, then evaluate a function or operation indicated by the user, returning the result.

The runtime name/value pair environment is held by the Context class and its sub-classes. A Context extends a `HashMap<LexNameToken, Value>`, so it is capable of storing and retrieving named values. In the same way that Environment objects were chained together during type checking. Context objects can be chained together as the evaluation creates new named values, and those names later disappear from scope.

Contexts allow named values to be retrieved by searching the present context, and then subsequent chained contexts. Hence a context for global variables might be chained with one for a function's parameters, and a further one defining variables in a "let" expression. Then a lookup of a variable would search this chain in reverse order.

In the case of functions or operations calling other functions or operations, the name searching should not proceed down the context chain beyond the nearest function or operation point – ie. if func1 calls func2, func1's variables are not in scope in func2. But global variables are visible in this case. Therefore a sub-class of Context, called RootContext is used to represent points in the context chain where the search should "jump" down to the global level. In fact there are three sub-classes of RootContext, one of which, ObjectContext, is specialized to hold an object value for "self" (which is in scope at the start of object member calls), another, ClassContext refers to a ClassDefinition for static invocations, and the third, StateContext, is able to have a module's state data attached (the actual state that is visible in VDM-SL changes as the operation call stack jumps from module to module, so this must be held in the context chain somewhere).

To create the global environment, the interpreter asks the default module or the ClassList passed to create the name/values from their definitions. The ClassList initialize method will dump all class' public static values into one global public static Context object; while the ModuleList initialize method will cause each module to initialize itself. The ClassInterpreter will then take the global public static scope and add the private static scope of the default class before starting. The ModuleInterpreter just takes the initial context from the default module before starting.

The code for the two interpreters' execute methods are similar therefore. Here is ClassInterpreter's:

```
@Override
public Value execute(String line, DBGPRReader dbgp) throws Exception
{
    Expression expr = parseExpression(line, getDefaultName());
    Environment env = getGlobalEnvironment();
    Environment created = new FlatCheckedEnvironment(
        createdDefinitions.asList(), env, NameScope.NAMESANDSTATE);

    typeCheck(expr, created);
    return execute(expr, dbgp);
}

private Value execute(Expression expr, DBGPRReader dbgp) throws Exception
{
    Context mainContext = new StateContext(
        defaultClass.name.location, "global static scope");

    mainContext.putAll(initialContext);
    mainContext.putAll(createdValues);
    mainContext.setThreadState(dbgp, CPUValue.vCPU);
    clearBreakpointHits();

    scheduler.reset();
    MainThread main = new MainThread(expr, mainContext);
    main.start();
    scheduler.start(main);

    return main.getResult(); // Can throw ContextException
}
```

Note that the public method is given a string expression to evaluate, so first it parses and type checks the expression given. The private method contains the important bit: a new StateContext is created,

the global public static content is added, any user-created values are added, the thread state is initialized to reference the virtual CPU (for VDM-RT), the breakpoint hit counts are cleared, and lastly the expression passed is evaluated in the context constructed by creating a MainThread (see the scheduler section below) and using the interpreter's scheduler instance to coordinate any further threads that the evaluation may create. The return value is finally extracted from the completed main thread.

When processing reaches a breakpoint, the evaluation of the current expression or statement should suspend, and enter a state where the interactive user can examine the state of the system.

This is fairly simple with VDM-SL because the specification cannot have threads, therefore pausing at a breakpoint must pause the entire system, and the user can easily examine anything they wish to. However, with VDM++ and VDM-RT the presence of threads make debugging more complex. It can be difficult to debug multiple threads from a single-threaded console.

The problem of single threaded consoles is conveniently sidestepped by using a GUI debugger, or more precisely, using the DBGp protocol to debug a remote VDMJ process. The advantage is that the protocol creates a separate TCP connection to the GUI for every thread that is created. This then avoids the problem of several threads wanting control of "the terminal". A command line DBGp client was provided for *real men* who prefer such things. Called VDMJC, this switches the console between connected DBGp threads by user command.

All breakpoint handling in VDMJ uses a Breakpoint class, and its subclasses. Every Statement and Expression has a breakpoint member object, and the eval methods of Expressions and Statements call their breakpoint's check method:

```
public abstract class Expression implements Serializable
{
    /** The textual location of the expression. */
    public final LexLocation location;

    /** The expression's breakpoint, if any. */
    public Breakpoint breakpoint;
```

and in subclasses ...

```
@Override
public Value eval(Context ctxt)
{
    breakpoint.check(location, ctxt);
    ...
```

The constructor for Expression and Statement just set the breakpoint field to a default Breakpoint. This base class' check method will not do anything under normal circumstances, but when a breakpoint is created by the user, the object in the expression or statement concerned is *replaced* with a subclass object, called Stoppoint. The check method of a Stoppoint will stop, although it may evaluate a conditional expression first.

When a Stoppoint stops, it allows the user to interact with the system at that point, evaluate expressions, print the stack and so on. This is done by calling some sort of command interpreter to read the console, or by calling into the DBGp session code to communicate with the GUI. However this is done, the important point is that when the user asks the thread to resume, control returns into the check method, and then out into the eval method of the stopped expression or statement.

If the user has issued a "step" or "next" or "out" command, we want to stop at the next line (skipping calls with "next", or until the current function/operation returns with "out"). This means that the Breakpoint base classes encountered as the execution proceeds have to check this – the system does not set any new breakpoints when stepping as it does not know where the code will go next. Therefore the check code for base class Breakpoint has to check the debugging state – and do so as efficiently as possible as this is the default code that is always called.

A simplified version of the Stoppoint version of check is as follows:

```
public void check(LexLocation execl, Context ctxt)
{
    location.hit();
    hits++;

    if (condition == null || condition.eval(ctxt))
    {
        SchedulableThread.suspendOthers();

        if (Settings.usingDBGP)
        {
            ctxt.threadState.dbgp.stopped(ctxt, this);
        }
        else
        {
            new DebuggerReader(ctxt, this).run();
        }
    }
}
```

The hit counters are to do with conditional breakpoint hits, and code coverage. Note that, because the check method is called for every Expression and Statement that is executed, it is a good place to record code coverage (against the location).

Any conditional expression will have been parsed (ie. done once) when the Stoppoint was created, and a test simply evaluates any boolean expression using the Context passed. Note that if the thread is about to stop, it also suspends the other threads. This is so that their state can also be examined by the debugger. After that, depending on whether we are using the command line or the GUI, the method calls one of two interactive methods; these return when the user continues.

The check method of the base class Breakpoint has to handle step, next and out rather than conditional breakpoints. Notice that step will stop whenever the location is not equal to the one in the state – ie. we are on a different file/line (but not a different position on the same line):

```
public void check(LexLocation execl, Context ctxt)
{
    location.hit();
    hits++;

    ThreadState state = ctxt.threadState;

    if (Settings.dialect != Dialect.VDM_SL)
    {
        state.reschedule();
    }

    if (state.stepline != null)
    {
        if (!execl.equals(state.stepline))
        {
            if ((stepping) ||
                (next && !isAboveNext(ctxt.getRoot())) ||
                (out && isOutOrBelow(ctxt)))
            {
                new Stoppoint().check(location, ctxt);
            }
        }
    }
}
```

As well as marking the code coverage and breakpoint hit counter, the base Breakpoint's check

method also calls the reschedule method of ThreadState which will conditionally terminate a timeslice if sufficient statements or expressions have been evaluated (see SchedulableThread below). Normally, this method just returns.

The isAboveNext and isOutOrBelow methods look at the context stack to decide whether to stop based on whether we are in a deeper stack frame (skip for next) or a shallower one (stop for out).

Notice that to actually stop, we create a Stoppoint and call its check method (which stops unconditionally), so it behaves as though you had inserted a breakpoint at the current location.

There is a third type of breakpoint called a Tracepoint. This is very simple as it does not interrupt the flow of control in its check method, but rather evaluates some trace expression and prints its value to the console (or sends it to the GUI).

2.14.1. Comments

The initialization of the VDM++ and VDM-RT runtime system is quite complicated, and should probably be cleaned up.

Here are a few points of general interest:

- The runtime system is initialized from the ClassList or ModuleList classes' initialize method.
- VDM++ and VDM-RT have to initialize the statics for each class first. This is done in two steps via methods on the ClassDefinitions. The first, staticInit creates the functions, operations and types – these do not have initializers; the second staticValuesInit covers instance variables and values, which can make calls to other static methods in their initializers.
- Value initializers can make forward references to other values which are not yet initialized. This causes characteristic exceptions, which are caught (ignored) and produce a second pass of the staticValuesInit initialization – up to a limit (currently 3). This ought to be done by working out the reference graph for the initializers.
- The static data thus initialized is written to a Context passed in (which becomes the global static environment). Static name value pairs are also written to maps inside the ClassDefinition – static data resides "in" the ClassDefinition at runtime.

And regarding object creation:

- The construction of an object (an ObjectValue) creates the superclass objects first, then for inherited fields, it adds a "locally named" **reference** to the superclass value. This is so that they can be explicitly referred to with a local name like C`x at runtime.
- Lastly, local definitions are used to override anything inherited, and the set of members are passed, together with the superclass objects, to the ObjectValue constructor.
- If the class defines an operation with the same name as the class (a constructor), it is called.

At runtime, an object's "self" field values are accessed via an ObjectContext which refers to the self ObjectValue. The check method of the context first searches the local contents for the name, then uses the "get" method of the ObjectValue (and lastly uses globals, if the name is not in self). The ObjectValue's get method is told whether the name is explicit or implicit (ie. whether it is qualified with a class name and a backtick), and uses this to either create a local name and search each object in the contained hierarchy, or uses the explicit name to search the hierarchy. If a static function or operation is called, there is no self, only static values. These are managed via a ClassContext, which refers to a ClassDefinition rather than an ObjectValue. The ClassDefinition's "get" method is conceptually the same as the ObjectValue's except it searches the static values of the class hierarchy.

2.15. vdmj.scheduler

Class Summary	
ResourceScheduler	The master resource scheduler.
SystemClock	The system clock, used by VDM++ and VDM-RT.
Resource	An abstract resource, to be scheduled.
BUSResource	A BUS resource.
CPUResource	A CPU resource.
SchedulingPolicy	An abstract scheduling policy.
FCFSPolicy	The First Come First Served scheduling policy.
FPPolicy	The Fixed Priority scheduling policy.
SchedulableThread	An abstract thread that can be scheduled by the resource scheduler.
MainThread	The main thread of an evaluation
ObjectThread	A thread created by an object start statement (VDM++)
AsyncThread	An asynchronous thread (VDM-RT)
BusThread	A thread handling a BUS (VDM-RT)
PeriodicThread	A periodic thread (VDM++ or VDM-RT)
CTMainThread	A main thread for running combinatorial tests
ControlQueue	An exclusive lock for a resource
Holder<T>	A pool for synchronized passing of data
Lock	A lock for handling guards
MessagePacket	An abstract message packet
MessageRequest	An operation request packet
MessageResponse	An operation reply packet.

Enum Summary	
Signal	An instruction to a suspended thread.
RunState	A SchedulableThread's current run state.

The vdmj.scheduler package defines classes that schedule the deterministic execution of multiple threads, and for VDM-RT, coordinate the movement of simulated time.

VDM-SL specifications are simple single threaded evaluations. Their execution will always produce the same result (unless they use a random input). However, VDM++ and VDM-RT specifications can have multiple threads of execution, and potentially the evaluation becomes non-deterministic. In order to prevent this, VDMJ uses a scheduler which coordinates the activity of all threads in the system and allows them to proceed, according to a policy, in a deterministic order. This guarantees repeatable evaluations even for highly threaded specifications.

VDMJ implements VDM threads using Java threads. The Java language does not define the operation of the underlying JVM thread scheduler. Instead, Java places various constraints on the ordering of events that must occur between threads, in the light of synchronization primitives that control access to shared resources. This means that although a given Java program will have a well defined partial ordering of some of its operations, the JVM thread scheduler has a great deal of freedom regarding how to execute threads that are not using synchronized access to variables or methods. In particular,

there is no way to control which thread gets control of which (real) CPU in the system, or for how long, though this can be influenced by setting a thread priority.

Building on this, VDMJ has to use Java synchronization primitives to enforce the semantics of the various VDM language features to control concurrency (permission guards and mutexes), and to control the order and duration of timeslices allocated to the different threads.

VDMJ scheduling is controlled on the basis of multiple "Resources" by a "Resource Scheduler". A resource is a separate limited resource in the system, such as a CPU or a BUS. These are separate in the sense that multiple CPUs or BUSES may exist, and limited in the sense that one CPU can only run one thread at a time, and one BUS can only be transmitting one message at a time. Therefore there is a queue of activity that should be scheduled for each resource – threads to run on a CPU, or messages to be sent via a BUS. The Resource Scheduler is responsible for scheduling access to the resources in the system.

The ResourceScheduler class implements the master resource scheduler. One interpreter (of any dialect) has a single ResourceScheduler instance. The ResourceScheduler controls a list of Resource objects, each of which is added to the list via a register method. The abstract Resource class has concrete subclasses called CPUResource and BUSResource. A VDM-SL or VDM++ system will have only one CPUResource (called a *virtual* CPU) and no BUSResources; a VDM-RT system will have as many CPUs and BUSES as are defined in the "system" class.

Every Resource has a scheduling policy, potentially different for each instance of the resource. A policy extends the abstract SchedulingPolicy class, and implements methods to reschedule (to reconsider what is best to run next) and subsequently to identify the SchedulableThread (below) that is best to run next and for how long it is to be allowed to run (a timeslice). With VDM-RT, in the event that all of its threads are blocked waiting to move system time, the policy will identify this fact and identify the minimum time step that would satisfy at least one of the waiters.

The ResourceScheduler loops through its resources with the following method outline:

```
do
{
    long minstep = Long.MAX_VALUE;
    idle = true;

    for (Resource resource: resources)
    {
        if (resource.reschedule())
        {
            idle = false;
        }
        else
        {
            long d = resource.getMinimumTimestep();

            if (d < minstep)
            {
                minstep = d;
            }
        }
    }

    if (idle && minstep has been set)
    {
        SystemClock.advance(minstep);

        for (Resource resource: resources)
        {
            resource.advance();
        }

        idle = false;
    }
}
```

```
while (!idle && main thread is not COMPLETE);
```

The Resource's reschedule method is responsible for determining the best thread to run next, if any, and to run it for the policy-determined timeslice. If the reschedule actually did anything (ie. the resource is not idle), it returns true. If it is idle, it is possible that the resource has all its runnable threads waiting to perform a time step (for VDM-RT). In that case, the getMinimumTimestep method returns that value, and the scheduler uses it to find the global minimum time step.

If all Resources have been considered, and the system is completely idle, and the global minimum time step has actually been set (the resources may be idle because they are waiting on guards or messages rather than moving time), then global system time can be moved by the global minimum step. This is done via a call to SystemClock, which manages global time, followed by a call to each resource's advance method, which delegates to the policy's advance method. This will result in all threads which were previously in the TIMESTEP state (blocked trying to move time) being moved to the RUNNABLE state (these are RunState enum values), and any which are now ready to run will have their transaction variables committed. The scheduler then notes that it is not idle, and goes back to reschedule all the resources. One of the (now) runnable threads will be scheduled, according to policy.

The reschedule method for the CPUResource is the most complex of the two resource types. It performs the following outline:

```
if (policy.reschedule())
{
    SchedulableThread best = policy.getThread();

    if (running != best)    // We swapped
    {
        if (running != null)
        {
            RTLogger.log("ThreadSwapOut...");
        }

        if (delayed)
        {
            RTLogger.log("DelayedThreadSwapIn...");
        }
        else
        {
            RTLogger.log("ThreadSwapIn...");
        }
    }

    running = best;
    running.runslice(policy.getTimeslice());

    if (running.getRunState() == RunState.COMPLETE)
    {
        RTLogger.log("ThreadSwapOut...");
        RTLogger.log("ThreadKill...");
        running = null;
    }

    return true;
}
else
{
    return false;
}
```

This first uses the policy to determine the best thread to run next. If there is none (the CPU is idle) it returns false. Otherwise the best thread is picked up, and if it is not the one currently running on the

CPU, the old one is swapped out and the new one is swapped in (possibly with a delay).

The new running thread is then set, and it is allowed to run for the timeslice given by the policy. Then when the timeslice is over (or the execution ends) if the thread is complete the thread is swapped out and killed. In any event, true is returned to the resource scheduler because this resource did something (it was not idle).

There are two policies implemented: FCFSPolicy and FPPolicy. They both extend SchedulingPolicy. These are actually the classes that maintain the list of threads allocated to a Resource (every Resource has a policy). The only difference between them is in the way they decide on a timeslice. The FCFS policy returns a fixed timeslice (ie. the execution of a fixed number of operations or expressions is permitted), whereas the FP policy returns a variable number, actually being identical to the number passed to setPriority (in VDM-RT). If a thread has no priority set, it defaults to the FCFS timeslice. This means that both policies give every thread a chance to run, but the FP policy gives higher priority threads a longer time on the CPU when they do run.

Both policies can have a certain amount of "jitter" added to their calculation of timeslices. This is set via the VDMJ properties file (see section 2.20). There is no jitter by default, which produces perfectly deterministic evaluations every time. However, this can hide problems with VDM thread synchronization that just happen to work because of the scheduling. Adding jitter will cause the thread scheduling to become slightly non-deterministic and therefore may expose such problems.

Back in the resource scheduler, if every thread is idle (and we cannot perform a time step), then one of two things has happened: the specification has completed because the main thread is complete, or we are deadlocked. The deadlock state is detected if the main thread is not complete and yet one of the resources has an "active" thread – ie. something which would be expected to proceed eventually, in state TIMESTEP or WAITING. If the specification is deadlocked, all threads are sent a DEADLOCKED signal, which trips them into the debugger. We wait for the main thread to end (from the debugger perhaps) before cleaning up.

```
if (main is not COMPLETE)
{
    for (Resource resource: resources)
    {
        if (resource.hasActive())
        {
            print("DEADLOCK detected");
            SchedulableThread.signalAll(Signal.DEADLOCKED);
            // Wait for main to terminate...
            break;
        }
    }
}

SchedulableThread.signalAll(Signal.TERMINATE);
```

Note that if the main thread is COMPLETE, we don't care about the state of the other resources. At the end of the execution, all threads are sent a TERMINATE signal, which causes them to throw a ThreadDeath exception (which kills a Java thread silently).

All thread resources that are scheduled by the resource scheduler via the Resource subclasses extend SchedulableThread (which in turn extends java.lang.Thread). The SchedulableThread's run method (ie. the body of the thread) is as follows:

```
@Override
public void run()
{
    reschedule();
    body();
    setState(COMPLETE);
    resource.unregister(this);
}
```

```
        synchronized (allThreads)
        {
            allThreads.remove(this);
        }
    }

    abstract protected void body();
```

The reschedule call basically turns the threads from state CREATED to RUNNABLE. Note that only one thread is ever RUNNING, but many can be RUNNABLE. The body is an abstract method that subclasses must provide, and after running the body, the state is changed to COMPLETE and the thread is unregistered from its resource and removed from the (static) list of all known threads.

The main resource scheduler and all the sub-resource scheduling calls run in the main Java thread. So as we've seen in the CPUResource reschedule above, a SchedulableThread is given a timeslice to run by the main thread calling the SchedulableThread's runslice method. That is responsible for putting the thread in the RUNNING state, running it for the given slice or until it changes its own state to something other than RUNNING, and then signalling the scheduler that it is done by returning from runslice.

This inter-thread coordination can only be achieved with Java synchronization primitives. The runslice method in SchedulableThread is as follows:

```
public synchronized void runslice(long slice)
{
    // Run one time slice - called by Scheduler
    timeslice = slice;
    waitWhileState(RunState.RUNNING, RunState.RUNNING);
}
```

The waitWhileState method puts the thread into the state given by the first argument, and then waits on the SchedulableThread's lock while the thread is in the second argument state – ie. set this thread as running, and don't come back until it is no longer running. There is a similar method called waitUntilState:

```
private synchronized void waitWhileState(
    RunState newstate, RunState until)
{
    setState(newstate);    // Call Java's notifyAll

    while (state == until)
    {
        sleep();           // Call Java's wait
    }
}

private synchronized void waitUntilState(
    RunState newstate, RunState until)
{
    setState(newstate);    // Call Java's notifyAll

    while (state != until)
    {
        sleep();           // Call Java's wait
    }
}
```

The act of changing the state does a notifyAll on the SchedulableThread's lock, so anything else that is calling sleep() will wake up and check the new state. So any thread calling waitUntilState that is waiting for it to become RUNNING will continue, and the scheduler will sleep until something else changes the state to something other than RUNNING, such as the thread calling the waiting method:

```
public synchronized void waiting()
{
    // Enter a waiting state - called by thread
    waitUntilState(WAITING, RUNNING);
}
```

This would be called by the `SchedulableThread` code when it decides that it needs to suspend in the `WAITING` state, such as when it is waiting for a message on a BUS. Similarly `SchedulableThread`'s `reschedule` method is called when the timeslice runs out to keep the thread runnable, but yield control:

```
private synchronized void reschedule()
{
    // Yield control but remain runnable - called by thread
    waitUntilState(RUNNABLE, RUNNING);
}
```

Notice that all of these methods are synchronized on the `SchedulableThread`'s lock. This is so that they can receive the `notifyAll`, as well as because they depend on shared values changed by other threads.

The `sleep` method can be woken up by a "signal" as well as a "notifyAll" state change. Signals are used to trigger the same behaviour in all threads (typically), such as when a deadlock is reached or when the evaluation terminates. The signal is sent via the (Java) thread's `interrupt` method rather than `notifyAll`. This causes any blocked wait calls to throw an `InterruptedException`. So the `sleep` method is as follows:

```
private synchronized void sleep()
{
    while (true)
    {
        try
        {
            wait();      // For a state change notify
            return;
        }
        catch (InterruptedException e)
        {
            handleSignal(signal);
        }
    }
}

private synchronized void setSignal(Signal sig)
{
    signal = sig;
    interrupt();
}
```

Depending on the signal set, the `handleSignal` method may kill the thread, or trap into the debugger (at a breakpoint or a deadlock). After leaving the debugger (ie. after a "continue" or "step") the loop goes back to wait for a state change. In this way, the debugger cannot change the scheduled flow of control, though the behaviour as perceived by the debugging user can be rather "strange".

These primitives for controlling scheduling are built into more complex calls to achieve the effects we want. For example, when a VDM-RT thread wants to perform a "duration" time step, it has to suspend and wait for the system to send through an "advance" call when time has changed, then it needs to find out whether it got as much time as it needed (or whether a smaller requirement was satisfied for

another thread), suspending again with a smaller request if not. The support for this is in the duration method of SchedulableThread:

```
public synchronized void duration(long pause)
{
    // Wait until pause has passed - called by thread

    if (!inOuterTimeStep)
    {
        setTimestep(pause);
        durationEnd = SystemClock.getWallTime() + pause;

        do
        {
            waitUntilState(RunState.TIMESTEP, RunState.RUNNING);
            setTimestep(durationEnd - SystemClock.getWallTime());
        }
        while (getTimestep() > 0);

        setTimestep(Long.MAX_VALUE); // Finished
    }
}
```

The calls to setTimestep just set a value which is subsequently retrieved by the scheduler via the CPU resource when it is asking for the minimum time step. (Note that this is not performed if an "outer time step" is in progress – such as when a cycles or duration statement is timing a block of statements) The system time is checked in a loop that waits in state TIMESTEP until the time required (possibly in several smaller steps, if other threads want smaller amounts of time). Note that a pause of zero will still cause the thread to do one time step.

The duration call is used by CycleStatements and DurationStatements, as well as PeriodicThreads (below), but the most frequent use is by the simple execution of statements and expressions, which call the SchedulableThread's step method:

```
public void step()
{
    if (dialect is VDM_RT)
    {
        if (!virtual) // vCPUs don't take any time
        {
            duration(Properties.rt_duration_default);
        }
    }
    else
    {
        SystemClock.advance(Properties.rt_duration_default);
    }

    if (++steps >= timeslice)
    {
        reschedule();
        steps = 0;
    }
}
```

Note that VDM-RT calls duration; VDM++ simply advances the clock as it does not coordinate time between threads. If the timeslice for the thread has expired, reschedule yields control to the resource scheduler while keeping the thread RUNNABLE.

Most of the subclasses of `SchedulableThread` do simple things that involve evaluating an expression and catching any thrown exceptions. The expression to run is passed to the thread constructor. In the case of `AsyncThreads`, this is passed via a `MessageRequest` that has come via a BUS and the result is returned on the same BUS via a `MessageResponse`. The client of an async thread (assuming there is one – this is the inter-CPU case, rather than an operation marked as "async") waits for the result to be placed in a `Holder<T>` (here, a `Holder<MessageResponse>`). This is simply a data area of the `<T>` type given, plus a `ControlQueue` to wait on. A `ControlQueue` is something which a thread gains exclusive access to (via `join/leave` methods), with `block/stim` methods to wait for/signal an event. These are used to protect BUS resources and for a BUS reply to be signalled back to the client waiting. The `block` method of a `ControlQueue` calls the `waiting` method of its `SchedulableThread`. A `Lock` is a similar coordination device, but this is used during the evaluation of guards on synchronized operations. In this case, we want all threads trying to acquire the lock to be able to queue for it (and wait), but for all of them to wake up when the guard value has changed. This is very similar to a Java synchronization lock, except implemented in terms of the resource scheduler primitives.

A `PeriodicThread` is slightly more interesting than the others. It has to perform its operation periodically, but it also has to do so under strict time control, with various options for constraining the behaviour. This is enabled (for VDM-RT) by calling the `SchedulableThread`'s `duration` method of course:

```
protected void body()
{
    if (first time)
    {
        if (offset > 0 || jitter > 0)
        {
            waitUntil(offset + noise);
        }
    }
    else
    {
        waitUntil(expected);
    }

    new PeriodicThread(exected from nextTime()).start();

    run operation...
}

private void waitUntil(long until)
{
    long time = SystemClock.getWallTime();

    if (dialect is VDM_RT)
    {
        if (until > time)
        {
            duration(until - time);
        }
    }
    else
    {
        while (until > time)
        {
            reschedule();
            time = SystemClock.getWallTime();
        }
    }
}
```

Here the `waitUntil` method calls `duration` for the time remaining. Note that the VDM++ dialect does not use the `duration` method. This is because VDM++ specifications do not use coordinated time

movement – the system clock is incremented unconditionally by VDM++ in statements and expressions, rather than making coordinated time steps with other threads. This means that the periodic simply has to wait (by calling `reschedule`) until the time has moved beyond the point required.

The `waitUntil` method is called at the start of the thread, which either waits for the initial offset time or until its expected time before (a) creating a new thread with the *next* expected time, and (b) running the operation. The next thread waits until its expected time, and so on (which means that periodic iterations can overlap if the operation takes longer than the period). The `nextTime` method performs a calculation based on jitter and delays, using a PRNG (remember a PRNG `nextLong` call can be positive or negative):

```
private long nextTime()
{
    // "expected" was last run time, the next is one "period" away,
    // but this is influenced by jitter as long as it's at least
    // "delay" since "expected".

    long noise = (jitter == 0) ? 0 : PRNG.nextLong() % (jitter + 1);
    long next = SystemClock.getWallTime() + period + noise;

    if (delay > 0 && next - expected < delay) // Too close?
    {
        next = expected + delay;
    }

    return next;
}
```

Naively, the next time will just be "now" plus the period. But if there is a jitter value set, an amount of "noise" is calculated randomly between +jitter and -jitter and this is also added to the next time calculation. Lastly, if the resulting next time is too close to the previous one (less than the "delay" value) then the next value is set to precisely "delay" ticks after the last one was expected.

When `SchedulableThreads` try to invoke operations with a sync guard, they may need to enter a `WAITING` state until the guard value changes. This uses a `Lock` class as mentioned above. The method to test a sync guard is called "guard", and is called before invoking any operation that has a sync condition defined. The outline of the method is as follows:

```
private void guard() throws ValueException
{
    self.guardLock.lock();

    while (true)
    {
        synchronized (self) // So that test and act() are atomic
        {
            // We have to suspend thread swapping round the guard,
            // else we will reschedule another CPU thread while
            // having self locked, and that locks up everything!

            debug("guard TEST");
            setAtomic(true);
            boolean ok = guard.eval();
            setAtomic(false);

            if (ok)
            {
                debug("guard OK");
                act();
                break; // Out of while loop
            }
        }
    }
}
```

```

        // The guardLock list is signalled by the GuardValueListener
        // and by notifySelf when something changes.

        debug("guard WAIT");
        self.guardLock.block();
        debug("guard WAKE");
    }

    self.guardLock.unlock();
}

```

The important points are that the guardLock is held on a "self" basis – not per operation. That is because guards often refer to #act(op) etc for operations other than the current one, and so all changes to self should be signalled though a Lock held at that level. The other important point is that the test and the call to act() to increment #act must be made under a Java lock to avoid another thread sneaking in (though this is probably not necessary in the deterministic scheduler, it is still good practice – it was *essential* with the non-deterministic scheduler).

The self.guardLock is signalled (to wake up a blocked thread) by a GuardValueListener. This implements the ValueListener interface, and is associated with UpdatableValues that comprise the state which is read by each guard expression. Therefore changes to those values cause the guard to be re-evaluated. The guardLock is also signalled by the OperationValue's notifySelf method, which is called whenever any operation's #req, #act or #fin value changes – note that this is rather crude, unlike the UpdatableValues, and means that guards wake up for this reason more often than they need to.

Apart from BUSThread (which runs specialized code for BUS transfers) the other SchedulableThread types all evaluate some sort of expression in a context that they create at the start. This is done in their body methods. They also run one of two private methods to do this, runDBGP and runCmd, to handle the differences between a remote DBGP connection from the GUI and a local command line debugger. The two methods are like this, in outline:

```

private void runDBGP()
{
    try
    {
        // Do evaluation
    }
    catch (ValueException e)
    {
        suspendOthers();
        ResourceScheduler.setException(e);
        DBGPRReader.stopped(e);
    }
    catch (ContextException e)
    {
        suspendOthers();
        ResourceScheduler.setException(e);
        DBGPRReader.stopped(e);
    }
    catch (Exception e)
    {
        ResourceScheduler.setException(e);
        SchedulableThread.signalAll(SUSPEND);
    }
}

private void runCmd()
{
    try
    {
        // Do evaluation
    }
}

```

```

        catch (ValueException e)
        {
            suspendOthers();
            ResourceScheduler.setException(e);
            DebuggerReader.stopped(e);
        }
        catch (ContextException e)
        {
            suspendOthers();
            ResourceScheduler.setException(e);
            DebuggerReader.stopped(e);
        }
        catch (Exception e)
        {
            ResourceScheduler.setException(e);
            SchedulableThread.signalAll(Signal.SUSPEND);
        }
    }

```

The calls to setException on the resource scheduler cause the main thread to throw an exception at the very end of the evaluation (even if the exception was raised in a different thread). It also reports the exception on stderr. The catch clauses that enter the debugger call suspendOthers, which sends a signal to the other threads to suspend, trapping them into the debugger too (see above). For exceptions where we have no VDM context information, we just suspend all threads and let the current thread die (which we can't debug, as we have no context).

2.15.1. Comments

The amazing thing about this is that it actually performs quite well. :-)

The weakest part of the system is probably the way that history counter changes wake up all guards on the object rather than just those that refer to the history values concerned. The problem is that history values are not UpdatableValues and therefore cannot simply have GuardValueListeners attached. A related inefficiency is that, while a GuardValueListener will only trigger on value updates that are mentioned in a guard, it wakes up all guards on the object because the lock is object-wide.

This may cause problems with debugging because guard tests will be very frequently scheduled, and so control will flip to those threads very often, even when the guard cannot have changed.

The reschedule looping for VDM++ periodic threads (in waitUntil) is poor. It means that the threads are always ready to run, and will be scheduled (briefly) at regular intervals until their time occurs – though if timeslices are of the order of hundreds of steps, and a step is two ticks, and periods are typically a few thousand ticks apart, and there are several threads in the scheduling queue, perhaps the amount of unnecessary rescheduling is not actually that great.

The use of scheduler timeslice jitter (and the FP policy in general) does not have as great an affect with VDM-RT as one might imagine. This is because RT threads hardly ever exhaust their timeslice before suspending. Rather, thread suspension is usually caused by the duration that every statement has. This is not the case for VDM++ though.

2.16. vdmj.values

Interface Summary	
ValueListener	Implemented by classes that watch for changes to a Value.

Class Summary	
Value	The parent of all runtime values.
****Value	A value of type ****. There are about 30 of these.

NumericValue	The root of the numeric value types.
ReferenceValue	The root of the value classes which reference another value.
InvariantValue	A ReferenceValue which includes an invariant function.
UpdatableValue	A ReferenceValue in which the referenced value can be changed.
GuardValueListener	A listener for processing updates that affect operation guards
ClassInvariantListener	A listener for processing updates that affect object instance invariants.
NameValuePair	A class to hold a name and a runtime value pair.
NameValuePairList	A list of name/value pairs.
NameValuePairMap	A map of name/values.
Quantifier	A class representing a quantifier.
QuantifierSet	A class representing a set of quantifiers.
State	A class for holding a module's state data.
ValueList	A sequential list of values.
ValueMap	A map of value/values.
ValueSet	A set of values.
BUSValue	A VDM-RT interconnecting BUS.
CPUValue	A VDM-RT CPU.
TransactionValue	A VDM-RT thread-updated value.

The `vdmj.values` class contains a set of value classes to represent runtime values in the interpretation of a specification. They are all subclasses of the abstract `Value` class.

All `Values` implement a set of standard methods to allow them to work correctly with the Java collections framework: `equals`, `hashCode`, `toString`, `clone`.

The `convertTo` method takes a `Type` argument and is responsible for the dynamic type conversion of values from one type to another, where possible. This tests the `-dtc` settings flag, returning the value unchanged if the flag is set. The method is used when types are "enforced" in a specification, such as when values are passed as typed arguments, or when values are returned from a typed function or operation. The method is specialized in all the `Value` subclasses that generally know how to convert themselves into a small number of related types (eg. conversions between subclasses of `NumericValue`). If a value cannot convert itself, it calls up to its superclass. The `convertTo` at the root of the `Value` hierarchy deals with common complex cases: converting to a union (iteratively trying to convert to each type); converting to a parameter type (looking up the parameter's name to get its actual type); converting to optional types (convert to the underlying type); and converting to a named type (convert to the underlying type, then wrap with an `InvariantValue` to enforce any type invariant).

A `ReferenceValue` is an abstract class that contains a value, and delegates all operations to that contained value. The concrete subclasses are `InvariantValue`, which includes an invariant `FunctionValue` as well as the value; and `UpdatableValue`, which has a `set` method capable of changing the value referenced (and all its methods are synchronized to allow for safe access to the changed value from other threads). When an `UpdatableValue` is created, it is passed a list of `ValueListeners` (optionally), which are called back when the `set` method is called. This is used to invoke class or state invariant functions when values are changed, and to trigger guards in VDM++ and VDM-RT.

A `State` class holds a VDM-SL module's state, which is held in an `UpdatableValue` which contains a `RecordValue` – being the "mk_Sigma" record value that can be considered to hold the Sigma state. The class implements `ValueListener`, which means that any changes to the `UpdatableValues` which comprise the module state will automatically call the `changedValue` method of `State` to check whether any state invariant has been violated. The class includes a flag which suppresses the invariant check while the state is being initialized.

A `GuardValueListener` implements `ValueListener`, holding an `ObjectValue` to signal when it received a value update. This is used to identify when updates are made to variables which are identified in

operation permission guards in VDM++ and VDM-RT. The `changedValue` method calls the `signal` method of a `Lock` object in the `ObjectValue`, which is also waited for by operation guards.

A `ClassInvariantListener` also implements `ValueListener`, holding an invariant `OperationValue` which can be executed in the event that any of the dependent variables changes value. This is used to implement object instance invariants. The class includes a boolean flag, called `doInvariantChecks`, which can be used to disable checks – for example during a class construction, or during an atomic statement.

Most value classes are immutable, in the sense that they are constructed with an underlying value that is held in a (public) final field of the class, and never changes. The only exception to this is the `UpdatableValue`, whose referenced value can be modified by the `set` method – note this changes the immutable value referenced; it does not change the value of the referenced object.

`UpdatableValues` are used for “state” values (instance variables, module state and “dcl” values that can be changed in an assignment). They are often initialized with structured immutable values, so values implement a method called “`getUpdatable`” to create `UpdatableValue` versions of themselves and their sub-values.

A `Value` can be held together with a `LexNameToken` in a `NameValuePair`, and such pairs can be collected into a list or a map for convenience. For example, the members of an `ObjectValue` are held in a `NameValuePairMap`.

Internally, many value classes have basic Java types at their heart. The basic values underlying `SetValues`, `SeqValues` and `MapValues` are `ValueSets`, `ValueLists`, and `ValueMaps` respectively. The basic value of a `Value` can be extracted using one of several methods, like `realValue` or `functionValue`. If the `Value` concerned cannot be converted to the basic type sought, a `ValueException` is thrown. The basic underlying values are used in the `eval` method of expressions to (say) actually perform arithmetic.

To ensure deterministic behaviour of specifications, two value types must enforce an ordering on their contents where none is strictly necessary: sets and maps. If this was not done, expressions which iterate over these values could produce different results for different executions. Therefore, when a `SetValue` is constructed, its underlying `ValueSet` contents are sorted; similarly, a `MapValue` is based on a Java `TreeMap` which keeps the contents in a sorted order.

A `CPUValue` represents a CPU declaration in a VDM-RT system class. Creating such a value also creates a `CPUResource`, which registers itself with the resource scheduler. A `BUSValue` similarly represents a BUS connection in VDM-RT, creating and registering a `BUSResource` when it is constructed. A BUS comprises a link between a set of CPUs.

A `BUSValue` has two important methods: `transmit` and `reply`. These delegate though to the `BUSResource`, which creates an `AsyncThread` to handle a request, and signals a `Holder` with a reply.

In VDM-RT, when multiple threads access a shared variable – eg. a piece of state information in an object – updates to the variable value are only visible *outside* the current thread after the next timestep (while changes are obviously visible inside the modifying thread). A value cannot be modified by two threads concurrently in VDM-RT (this is a runtime error). This processing is managed by a `TransactionValue` class. This is a subclass of `UpdatableValue`, where the `set` method is overridden to write the new value to a local field rather than to the referenced value from the parent. The various retrieve methods (`intValue`, `booleanValue` etc) are overridden to use the new value for the thread that changed it, and the parent value otherwise. An additional `commit` method actually sets the parent value to the new value. The `commit` method is called during a time step, making the per-thread value visible to other threads.

Transactional variables are very new, and by default the commit system is disabled. It can be enabled by a property setting.

2.16.1. Comments

It seemed sensible to have all `Values` contain some sort of primitive value, including sets, sequences and maps, rather than having the `Value` classes extend the Java collection framework directly. That allows the `Value` types to be in a “VDM” hierarchy, but it leads to the confusing situation where a `MapValue` contains a `ValueMap` – the former extends `Value`, the latter extends

TreeMap<Value,Value>. The ValueXXX classes should only be used internally – ie. they should never be returned from an eval method (they can't be, as they're not Value subtypes).

There is a lot more discussion about VDM-RT and time handling in section 2.15.

2.17. vdmj.commands

Class Summary	
CommandReader	A class to read and perform commands from standard input.
ClassCommandReader	A class to read and perform class related commands from standard input.
ModuleCommandReader	A class to read and perform module related commands from standard input.
DebuggerReader	A class to read and perform debugging commands from standard input.

The vdmj.commands package contains the command line readers that implement the interactive actions of the suite. This should be the only package that interacts with the user terminal.

A subclass of CommandReader is instantiated to provide the main command prompt of the interpreter: either ClassCommandReader or ModuleCommandReader. These classes can set or clear breakpoints, but they do not permit commands like “step” and “next”, as these are only legal from a breakpoint context. The (few) differences between the two classes concern the commands that are legal for each – eg. “classes” and “modules” – while the truly common code is in the parent.

At a breakpoint, a DebuggerReader is used to permit the extra debugging commands allowed from this context, and prevent the ones that aren't (like initializing the environment on the fly). The one DebuggerReader is used to handle debugging sessions for VDM-SL, VDM++ and VDM-RT.

2.17.1. Comments

The idea here is to isolate the console interaction from the rest of the program so that the core interpreter could be included in (say) an Eclipse plugin. In practice, both Tracepoints and Stoppoints have to write to the console as well. This area needs some work to completely separate the implementation from the console control, so that breakpoints could interact with (say) Eclipse.

The CommandReader system works well for simple specifications, but multi-threaded debugging, especially in VDM-RT, is awkward. The *VDMJ Client*, which uses the DBGP protocol to a separate VDMJ process, is better for such debugging.

2.18. vdmj.messages

Class Summary	
VDMMessage	The root of all reported messages.
VDMError	A VDM error message.
VDMWarning	A VDM warning message.
Console	A class to provide System.in/out with a charset encoded wrapper.
Redirector	An abstract class to redirect output.
StdoutRedirector	A class to redirect standard output to a debugger.
StderrRedirector	A class to redirect standard error to a debugger.
RTLogger	A VDM-RT class for logging thread, CPU and BUS activity.

Exception Summary	
MessageException	An exception to carry a textual error message.
NumberedException	An exception to carry a number and text information.
LocatedException	An exception to carry number, text and location information.

The vdmj.messages package contains classes and exceptions to hold and display error and warning messages. Lists of VDMMessage values are returned from the TypeChecker and Reader classes. Similarly, internal exceptions that carry message numbers and position information are all subclasses of NumberedException.

The Console class manages output to stdout/stderr, in particular managing the character set to be used. The Redirector and its concrete subclasses are used to redirect or copy stdout and stderr to a debugger (see section 2.19).

The RTLogger is used by VDM-RT specifications to write loggable events – such that the timing diagram can be analysed using Overture's real-time log viewer. The class caches events in memory, occasionally flushing them out to a given PrintStream. By default this is the Console. It is also possible to turn logging off (the default state), which causes log events to be discarded.

2.19. vdmj.debug

Class Summary	
DBGPReader	The main class of the DBGp protocol reader.
DBGPCommand	A parsed IDE command.
DBGPOption	A parsed IDE command option.
DBGPFeatures	The set of features supported/set by the debugger.
RemoteControl	An interface for remote control of a VDM execution.
RemoteInterpreter	An interface for remote access to the interpreter.

Enum Summary	
DBGPBKbreakpointType	The possible DBGp breakpoint types.
DBGPCommandType	The possible DBGp IDE command types.
DBGPOptionType	The possible IDE command option flags.
DBGPErrCode	The possible status response error codes.
DBGPReason	The possible status response reason codes.
DBGPContextType	The possible variable name context types.
DBGPRedirect	The possible I/O redirection options.
DBGPStatus	The possible status responses.

Exception Summary	
DBGPEException	A general purpose debugger exception.

The vdmj.debug package contains classes that implement the DBGp remote debugging protocol defined in [6]. This allows VDMJ to load a set of specifications and evaluate/debug them under the remote control of another process – usually a GUI IDE, such as Eclipse.

As described in [6], the principle behind DBGp is that the IDE launches the debugged process (ie. VDMJ in our case), and that process connects back to the IDE using a TCP/IP connection to a host/port specified by the IDE. The IDE then sends commands to the debugger on the connection, and the debugger acts on those commands, sending status and results back to the IDE via the same connection.

Because DBGp starts VDMJ, the principal class that handles the connection, DBGPReader, defines a "main" method. This is separate from the main method defined in the VDMJ class (see section 2.1). The command line arguments are as follows:

```
Usage: -h <host> -p <port> -k <ide key> <-vdmpp|-vdmsl|-vdmrt>
      -e <expression> | -e64 <base64 expression>
      [-w] [-q] [-log <logfile URL>] [-c <charset>] [-r <release>]
      [-coverage <dir URL>] [-default64 <base64 name>]
      [-remote <class>] {<filename URLs>}
```

The host and port identify the connection that must be opened back to the IDE to receive commands. The ide key is a value that must be passed back to the IDE during the initial connection. The VDM dialect is a mandatory option. The expression, optionally base64 encoded for special charsets, is the main expression to be evaluated, and the list of filenames identify the specification itself. Each filename is in the form of a file URI ("[file:/...](#)"). There are also several optional settings. The -w and -q options suppress warnings and information messages respectively; the -log option sets the location of the VDM-RT real time log output; the -c option sets the character set for the specification; the -r option is a VDM language release name (classic or vdm10); the -coverage option defines where detailed coverage information is written (for use by the GUI); the -default64 option sets the default module or class name in base64; and the -remote option identifies a class name that will be called rather than the main VDMJ processing loop (for remote control – see section 3.2).

Depending on the dialect, DBGPReader creates an instance of VDMPP, VDMRT or VDMSL (see section 2.1), and uses it to parse and type check the list of files passed. The DBGp protocol assumes the "program" being debugged is already compiled correctly, so if there are any syntax or type checking errors, these are sent to the VDMJ process' standard output and the process quits with an error exit code. The protocol has no mechanism for returning these errors to the IDE.

If the specification is clean (warnings are permitted), an Interpreter object encapsulating the parsed/checked specification is obtained from the VDMSL, VDMPP or VDMRT object, and this, along with the host, port, ide key and expression are passed to the constructor of a new DBGPReader object. The constructor saves the values, and sends the DBGp initialization message back to the IDE. Lastly, the main method calls the run method of the reader. When this method returns, the VDMJ process quits with a success error code (0). If any exceptions are thrown from run, it quits with an error exit code.

The DBGPReader run method is a read/execute loop, reading DBGp commands from the IDE connection, processing them, sending back any responses, and then looping. The return value from the private methods to handle each type of command indicate whether the run method should keep looping or return (for example, the "detach" command would indicate that the loop should return, though most commands keep looping).

All DBGp commands are in a simple textural format, similar to the format of a UNIX command (see [6]):

```
command -op1 v1 -op2 v2 ... -- data
```

This text is parsed by a method which produces a DBGPCmd object; parsing errors throw a DBGPEException, which is caught and used to build an error status response before returning to the run loop.

Successfully parsed commands are passed to one of a number of processCmd methods which handle each type of command. The general form of these is to validate the options passed (held in the parsed DBGPCmd object), and send back an error response if not correct; and then act on the command, sending back a successful status, possibly combined with information being requested by

the IDE.

Responses sent to the IDE are all XML formatted messages (see [6]). A collection of private methods in DBGPRReader take the raw text response from a processed command in a StringBuilder, and use this to create an XML response message and send it on the connection to the IDE .

All of the DBGp commands which interact with the running specification do so by making method calls on the Interpreter object passed to the DBGPRReader's constructor. Note that this is the abstract Interpreter class, not the concrete ClassInterpreter or ModuleInterpreter, so the DBGPRReader will work with VDM-SL, VDM++ or VDM-RT.

When the specification is executed (via the DBGp "run" command), the interpreter's execute method is called, passing the expression to be evaluated. The DBGPRReader instance is also passed to the execute method; this is stored in the thread context information for the main VDM thread, and allows breakpoints to find the connection to the IDE.

The interpreter's execute method does not return until the specification has been fully evaluated, so this raises the question of what happens at breakpoints, when information must be passed to the IDE and more instructions received.

Breakpoints normally use the DebuggerReader class to interact with the user on the command line before continuing execution (see section 2.14). But when the process is being debugged by an IDE, the breakpoint will find the DBGPRReader object associated with its thread context, and call its "stopped" method. This sets the state of the connection to "break" and sends a status message to the IDE to let it know that execution has stopped at a breakpoint. It then enters the "run" method to process IDE commands as it did originally (in fact this is a recursive call, since the original run call is on the stack, having called the execute method). Note that some DBGp commands are only acceptable when the connection is in the "break" state, such as "step_into" and "stack_get". Expressions can be evaluated in the "break" state and will be evaluated in the context of the breakpoint (ie. local variables are visible).

A "continue" command from the IDE will cause the recursive run call to return, and the flow of control will return to the breakpoint and from there back into the main execution. Further breaks may occur, but eventually, the main evaluation will complete and the original run method call will regain control. Note that the debugged process does not automatically quit at this point. It enters a "stopped" state, where further (restricted) IDE commands are possible, such as to retrieve the final evaluation result.

The DBGp protocol allows multi-threaded programs to be debugged. In this case, each thread opens its own connection to the IDE (on the same host/port), though the opening of the connection is delayed until the thread has something to say – perhaps at a breakpoint – to avoid bombarding the GUI with thread create/death connections that do nothing. The IDE is responsible for sending commands for each thread on the appropriate connection. As far as the debugged process is concerned, these connections are completely separate. Therefore, when VDM starts a new thread, the creating thread's DBGPRReader is "cloned" to create a new object which will lazily open its own connection to the IDE, and be stored in the new thread's context. When the second thread reaches a breakpoint, it will respond to the IDE on its own connection.

The IDE can request that stdout/stderr output from the debugged process be sent to the IDE rather than (or as well as) to the usual console. If such a command is received, the Console class is called (see section 2.17), and a Redirector is added to the appropriate stream. These objects are passed a DBGPRReader reference and use this to send output to the IDE when directed to do so.

The RemoteControl and RemoteInterpreter classes are covered in section 3.

2.19.1. Comments

It would probably be better to use some sort of XML handling package to build the IDE responses, rather than hand-crafting them. Though the XML involved is very simple.

Asynchronous debugging could be done if checkpoints (ie. at the start of every expression or statement) check the status of the DBGp connection. That is simple to arrange, as the object is in the thread's context, but the overhead may be large (calling the input stream's "available" method). It might be acceptable to (say) only test the stream every 100 or 1000 operations, at the risk of making the threads unresponsive to asynchronous breaks.

A better solution would be to have asynchronous listening threads that call the signal method of `SchedulableThreads`, though that would double the number of threads in the system.

Note that the `Console` class only has one `Redirector` wrapped around an output stream. That means that all output is redirected to one particular DBGp connection, rather than the output from different threads being directed to each thread's DBGp connection. The thread that receives all the output is the one that last sent the IDE command to redirect it.

Currently, the values of all variables are sent back to the IDE as strings. To enable the structure of aggregate types to be expanded by the IDE, we need to define an XML Schema definition to describe them. This has not yet been done. When it is defined, it should be possible to change the `Type` class hierarchy to produce XML Schema to describe themselves, and to change the `Value` class hierarchy to "print" themselves in that schema.

2.20. vdmj.config

Class Summary	
Properties	A value to provide access to the <code>vdmj.properties</code> file.

This is a simple class with static fields, each representing (and updated with) values in the `vdmj.properties` file. The class extends `ConfigBase`, which contains reflection calls to lookup field names in the properties file (if any) and update the default static values with the property read. There is a `vdmj.properties` file included with VDMJ which just contains the default values:

```
#
# Settings for VDMJ. These override defaults in the code.
#

# The tab stop for source files.
# (default 4)
parser.tabstop = 4

# The default duration for RT statements
# (default 2)
rt.duration.default = 2

# The default timeslice (statements executed) for the FCFS policy
# (default 100)
scheduler.fcfs.timeslice = 100

# The vCPU/vBUS timeslice
# (default 10000)
scheduler.virtual.timeslice = 10000

# The timeslice variation (+/- jitter ticks)
# (default 0)
scheduler.jitter = 0

# Enable transactional variable updates
# (default false)
rt.duration.transactions = false

# Enable extra diagnostics for guards etc. */
# (default false)
diags.guards = false

# Enable extra RT log diagnostics for timesteps.
# (default false)
diags.timestep = false
```

2.21. vdmj.util

Class Summary	
Utils	A utility class.
Base64	A class to encode/decode base64 strings.
Delegate	A class to hold a native delegate Java object.
ConfigBase	A base class for property file readers.

Utils is a small utility class. It just defines a few methods for printing out a comma separated list of arbitrary types held in a Java List<T>.

Base64 does what you would expect. It is used by the DBGp protocol.

The Delegate class is used to lookup a Java class of the same name as a given VDM++ or VDM-RT class or VDM-SL module, and hold a reference to it. If such a class is found, and the VDMJ specification encounters an "is not yet specified" expression or statement, then the call is delegated to a similarly named method in the native delegate object. This is how the IO, MATH and VDMUtils classes/modules are implemented. See section 3.1.

ConfigBase supports the Property reading system, with reflection calls to examine a class of static fields, and update them with property values.

3. External Interfaces

3.1. The Native Call Interface

Parts of a VDM specification may wish to call out to native Java code, for example to compute mathematical library functions or to perform IO operations. The external calling mechanism in VDMJ uses the "is not yet specified" statement and expression to achieve this.

When a non-specified statement or expression is encountered, VDMJ determines whether the class or module containing it has a Java "delegate". A delegate is a Java object which has methods of the same name as the non-specified VDM operation or function, and to which the evaluation can be delegated, returning a real result rather than a runtime exception.

If there is no delegate, and this is the first time a non-specified statement or expression has been encountered in the current module or class, then VDMJ looks through the Java CLASSPATH for a Java class of the same name as the name of the module or class (with underscores turned into package separators). So for example, a module or class called `org_xyz_Handler` would look for a class in the `org.xyz` package called `Handler`. If such a Java class is found, an instance of the class is created, and associated with the Module, the `ClassDefinition` or the `ObjectValue`, depending on whether the operation or function being invoked is static.

The delegate object is thereafter associated with the module, class or object in VDM, and so state values written inside the Java object will be retained. Calls to the "is not yet implemented" operations or functions are passed to the delegate, with argument values passed as VDMJ Values. Similarly, the return value of the Java method must be a VDMJ Value.

It is a runtime error to have delegate methods which do not take Value arguments or return Value. It is also a runtime error to have non-static operations/functions delegate to static Java methods. You can only distinguish delegated overloaded names by the number of Value parameters (ie. you cannot distinguish overloaded members by the actual types of the parameters, as they are all Value in Java).

The delegate discovery call is only made once. Thereafter, if there is no delegate, the "is not yet specified" statement or expression will raise a normal `ContextException`.

Note that a class or module with a simple name like `MATH` or `IO` must have a delegate Java class in the default Java package. The `MATH`, `IO` and `VDMUtil` standard library functions in VDMJ are implemented using the same delegate scheme. It is possible to delegate from flat VDM-SL specifications, but the Java class has to be called `DEFAULT` (which is the VDM module name for flat specifications).

3.2. The Remote Control Interface

It is sometimes desirable to have VDM specifications "controlled" by external programs. To enable this, VDMJ provides a `RemoteControl` interface which must be implemented by external programs wishing to take control of a specification.

The interface only defines one method:

```
public interface RemoteControl
{
    public void run(RemoteInterpreter interpreter) throws Exception;
}
```

The class implementing this interface is identified to VDMJ using the `-remote` command line argument. If such an option is given, the specification is parsed and checked as normal, but instead of entering the interpreter, an instance of the remote class is instantiated (default constructor) and the `run` method is called, passing a `RemoteInterpreter` object. The class is responsible for implementing the `run` method such that control does not return until the program is complete. The `RemoteInterpreter` passed has three methods which can be called to interact with the specification: `execute`, `valueExecute` and

init. The first takes a string argument and evaluates it, returning a string value; the second takes a string argument and returns a VDMJ Value; the third re-initializes the interpreter.

Any exceptions (eg. runtime ContextExceptions) raised by the interpreter as a result of calling the methods on RemoteInterpreter will be returned to the controlling class.

One possible use of the remote control interface is with the JUnit testing framework to test VDM specifications. For example, a run method could be written as follows:

```
public void run(RemoteInterpreter i) throws Exception
{
    TestClass.init(i);
    TestRunner.run(TestClass.class);
}
```

This is a fairly crude way of invoking JUnit's text test runner, identifying a TestClass containing a set of tests which JUnit will iterate through, presumably calling the valueExecute method of the RemoteInterpreter, and asserting the value of the results or exceptions returned. The JUnit setUp method of TestClass can be used to initialize the interpreter between tests, if that is required.